

AutoDojo: Adaptive Black-Box Attacks Reveal the Limits of IPI Defenses and Task-Specification Effects in LLM Agents

Xinhang Ma*, Taoran Li†, Chaowei Xiao‡, Zhiyuan Yu†, Ning Zhang*, and Yevgeniy Vorobeychik*

*Washington University in St. Louis, St. Louis, MO, USA

†Texas A&M University, College Station, TX, USA

‡Johns Hopkins University, Baltimore, MD, USA

Abstract—Indirect prompt injection (IPI) is a major security threat to LLM-powered agents. Thus, a growing body of work have proposed a variety of defensive approaches against IPI. These can be grouped into three broad categories: 1) prompt-based (using prompting as a way to prevent agents from following malicious instructions), 2) detection-based (identifying and filtering malicious instructions), and 3) system-level (using systems insights, such as control and data isolation, for defense). However, commonly used benchmarks for evaluating defense, such as AgentDojo, are *inherently static*, generating a fixed distribution of IPI attacks. Consequently, static benchmarks do not usefully evaluate defense robustness to adaptive threats. We address this issue by developing AutoDojo, an adaptive extension of AgentDojo that optimizes IPI against a given defense. Using AutoDojo against state-of-the-art IPI defenses across three task suites and five target models, we make two key observations. First, many defenses offer only limited protection: a cheap, black-box adaptive attack using a frontier LLM to iteratively optimize the injection raises attack success rate (ASR) well above the level achieved by static injections against nearly all evaluated defenses. Against a filter that reduces static ASR to 0%, AutoDojo recovers 28% overall and 64% on action-open tasks. Second, for prompt-level and filter-based defenses, ASR is substantially higher on *action-open* tasks—where the user’s request delegates the action itself to attacker-controlled content—than on precisely specified tasks. This is a structural limit: on such tasks the injection can pose as ordinary data rather than an explicit instruction, bypassing defenses that rely on detecting instruction-like text. AutoDojo is publicly available at <https://github.com/xhOwenMa/AutoDojo>.

1. Introduction

LLM agents increasingly act on untrusted external data, and this has made *indirect prompt injection* (IPI) a central security concern for tool-using systems [1]. In an IPI attack, adversarial instructions embedded in tool outputs, retrieved documents, or third-party content hijack the agent away from the user’s intended task. The risk is amplified because an agent acts through consequential and often irreversible tools, such as paying a bill, sending a message, or sharing a document, so a single hijacked step can cause real harm.

Moreover, the attack surface grows with every external source the agent is permitted to read.

A large and fast-growing body of work proposes defenses against IPI. These include classifiers that filter injected text [2], [3], [4], prompt-level schemes that separate instructions from data [5], [6], and methods that constrain the agent’s actions to its intended task [7], [8], [9], [10]; we refer the reader to recent surveys for a fuller treatment of this fast-moving landscape [11]. These defenses report strong protection, but nearly all are evaluated the same way: against the fixed, hand-written injection strings that benchmarks such as AgentDojo [12] ship with, summarized as a single aggregate attack success rate (ASR). Such a number measures robustness to a *frozen* attack, not to an adversary who adapts. Consequently, a defense can score well on it by keying on superficial features of the canonical strings rather than on any robust notion of malicious intent.

Adaptive attacks, however, are not hypothetical. Paraphrasing, black-box search, and reinforcement-learned attackers already surpass hand-written injections in the broader prompt-injection literature [13], [14], [15]. Recent work pushes this to its limit, showing that by scaling general-purpose optimizers under a white-box attacker with full knowledge of the defense, many published defenses can be driven back to high attack success [16]. This is a valuable worst-case stress test, but its adversary is largely theoretical, far stronger than anyone who can merely query a deployed agent. The question that matters for practice is under what conditions a *realistic*, resource-light attacker suffices.

To answer this question, we develop *AutoDojo*, an evaluation framework that extends AgentDojo [12] by turning any existing injection into an adaptive one against a target LLM agent. The threat model in AutoDojo differs from the white-box worst case in three ways. First, it is strictly *black-box*, reflecting more realistic attack scenarios, and using only the success/fail signal that any external user can observe, with no model weights or gradients. Second, it is *cheap*: it optimizes an injection using a small query budget to an offline optimizer LLM, rather than performing a large-scale search or attacker model fine-tuning. Furthermore, AutoDojo naturally supports an evolving attack landscape by using any attacks as seeds in its optimization procedure.

We use AutoDojo to evaluate three broad classes of defenses: prompt-level schemes, detection and filtering-based




Figure 1. **Published IPI defenses are often only superficially robust.** For five defenses that work very well against the static `important_instructions` attack on a GPT-4o-mini agent (hollow markers, near the floor), AutoDojo—our cheap, black-box adaptive attack—recovers attack success (solid markers) against most of them. Attack success rate is aggregated over three task suites—banking, slack, travel—of the popular AgentDojo [12] benchmark that our framework builds upon. This pattern holds across target models (see §5).

defenses, and system-level defenses, comprising 9 state-of-the-art defenses in total. Our experiments demonstrate that, across five target models and our diverse task suites, the proposed adaptive attack increases ASR for nearly all defensive approaches well above the level of the static injections, often by a wide margin (Figure 1 offers an illustration). In the sharpest case, a purpose-built injection classifier that almost entirely blocks the static attack is driven back to near its no-defense ASR. Nonetheless, we find that some defenses, especially system-level defenses, do provide some protection even in the face of this adaptive attack.

Additionally, we observe that effectiveness of attack against existing defenses is not spread evenly across tasks. Rather, attacks tend to be considerably more successful when the user’s own request is under-specified. Specifically, we group user tasks into three categories that represent decreasing levels of precision in task specification: 1) *fully-specified* (action and parameters given), 2) *param-open* (action given, parameters read from content), and 3) *action-open* (the user names no action and defers to whatever the content says). We find that attack success rates are much higher for tasks in categories 2 and 3 than for those in category 1 (fully specified). Notably, a typical use case of agents is likely to fall into the less precisely specified categories, as most users are likely not AI or security experts.

Summary of Contributions

In summary, we make the following contributions:

- An efficient (i.e., requiring very few iterations of an optimization routine and API calls to a capable LLM), black-box *adaptive evaluation framework* for IPI defenses that optimizes any seed attack against a defended agent using only feedback about attack success rate (ASR).

- A new categorization of task specification precision (fully specified, underspecified parameters, and underspecified action) to better understand structural sources of IPI risks.
- Extensive experimental evaluation of three classes of IPI defenses, showing that the efficacy of many published defenses appears to be greatly inflated due to the existing *static* evaluation paradigms. On the other hand, cheap black-box adaptive attacks significantly increase ASR.
- **AutoDojo**, an evaluation framework that bundles the optimization pipeline that adapts a seed attack to a defended agent, a library of pre-optimized injections across target models and defenses, and per-task underspecification labels, providing an easy-to-use tool for fine-grained adaptive adversarial evaluation of IPI defenses.

2. Related Work

2.1. Indirect Prompt Injection

Prompt injection was first studied in the *direct* setting, where a user supplies adversarial input that overrides the developer’s system prompt [17], and was later formalized and benchmarked as a general attack against LLM-integrated applications [18]. Greshake et al. [1] introduced the *indirect* variant (IPI), in which the adversary does not interact with the model at all but plants instructions inside data that the application later retrieves. When an agent reads that data into its context while completing a benign user’s request, it may treat the planted text as instructions and carry out attacker-chosen tool calls. The distinguishing feature, relative to a jailbreak, is that the user is benign and unaware: the adversary controls only the contents of some external surface (a hotel review, a calendar event, a shared document) that the agent legitimately needs to read.

2.2. Benchmarks for Agent Security

Several benchmarks measure how well LLM agents resist attacks, each pairing benign user tasks with attacker-controlled content and reporting an aggregate attack success rate. InjecAgent builds a suite of tool-use cases and splits attacker goals into direct harm to the user and exfiltration of private data [19]; BIPIA benchmarks indirect prompt injection and its defenses [20]; and Agent Security Bench (ASB) broadens the scope to ten application scenarios spanning attacks such as prompt injection and memory poisoning [21]. AgentDojo, which we build on, instead provides a dynamic tool-calling environment in which the agent runs a real task loop with tools rather than a single forward pass, letting us score a defense end to end rather than on isolated inputs [12]. These benchmarks share one design choice: the attacker content is fixed, so a defense is scored once against a hand-written injection rather than against an adversary that adapts to it. We keep the AgentDojo environment and add an attacker that adapts to the defense under a strict black-box setting.

2.3. IPI Defenses

Proposed defenses against indirect prompt injection fall into a small number of families, distinguished by where in the agent loop they intervene.

Filtering-based defenses screen the data the agent reads and drop or flag any input judged to be an injection. One line of work trains a dedicated classifier on the input text: Meta’s Prompt Guard and ProtectAI’s DeBERTa detector are cheap encoder classifiers [2], [22]. PiGuard targets the over-defense problem in which trigger words alone cause false positives [3], and other detectors classify injections directly from text embeddings [23]. A related line reads auxiliary signals rather than the raw text: attention patterns that flag injected spans [24], instruction-intent analysis that checks whether the agent is about to follow untrusted text [25], span localization that pinpoints the injection inside the input [26], and leave-one-out attribution that flags when an untrusted segment dominates a privileged action [27]. A second family replaces the classifier with an LLM sanitizer: DataFilter fine-tunes an LLM to strip adversarial spans while preserving benign content [4], and PromptArmor prompts an off-the-shelf LLM to remove injected instructions before the agent acts [28]. All of these act on the *surface form* of the data rather than on the agent’s actions.

Prompt-level defenses leave the data untouched but restructure the prompt so the model can tell instructions from data. Spotlighting [5] marks the provenance of untrusted spans with delimiters, datamarking, or encoding; the instruction hierarchy [6] trains the model to prioritize higher-privilege messages over third-party content; and simpler “sandwich” and “reminder” prompts re-assert the user’s request around the untrusted data. These defenses raise the bar for an injection to read as a legitimate instruction but do not change what counts as instruction-like text.

System-level defenses constrain what the agent may do rather than what it may read, and fall into four lines. The first trains the model to ignore instructions that arrive on the data channel [7], [8]. The second derives a permitted tool-call trajectory from the user’s request and blocks calls that deviate from it [9], [10]. The third governs how information flows through the agent. One group separates a trusted plan from untrusted data so the latter cannot alter control flow: CaMeL [29] extracts the control and data flow from the trusted query, f-secure [30] filters untrusted inputs out of a structured planning pipeline, and ACE [31] builds an abstract plan from trusted input and verifies by static analysis. Others track confidentiality and integrity labels through the computation to enforce flow policies [32], [33], isolate execution across components [34], or block privilege escalation inside the agent [35]. The fourth analyzes the execution trace after the fact: AgentArmor [36] builds a program-like graph from the trace and enforces policies on it, while AgentSentry [37] applies temporal analysis to locate injections. System-level defenses need not recognize the injection itself; they instead require a clean specification of the agent’s intended actions to constrain against.

2.4. Adaptive IPI Attacks

A growing body of work shows that the canonical, hand-written injection strings shipped with benchmarks badly understate attacker capability. Adaptive attacks tailor the injection to a specific target, and the agentic line inherits two mechanisms from the jailbreaking literature.

The first is *black-box iterative refinement* with an auxiliary LLM, in the lineage of PAIR [38], which uses an attacker LLM to repeatedly rewrite a prompt against a target it can only query. In the IPI setting, TopicAttack [13] uses an LLM to author a gradual topic transition into the injected instruction, minimizing the semantic gap to the surrounding context and sustaining high success even under deployed defenses; AgentVigil [14] performs black-box search over injection templates with Monte Carlo tree search; and RLHammer [15] fine-tunes an attacker LLM with reinforcement learning. These methods require only query access to the target but differ widely in cost. Our framework belongs to this PAIR-style family and sits at its cheapest end: a frontier LLM iterating from a seed under an attack-success-only signal, with no gradients, no reinforcement learning, and no fuzzing harness. We use it to show that even this minimal, black-box attacker collapses defenses that report strong robustness on static benchmarks.

The second is *gradient-based optimization*, which extends GCG [39] to indirect prompt injection: a momentum-enhanced gradient search constructs universal injection strings that transfer across inputs [40]. Neural Exec [41] learns reusable execution-trigger delimiters by differentiable search, and Imprompter [42] adapts GCG to the agentic setting, optimizing obfuscated prompts that coerce an agent into improper tool calls such as data exfiltration. Unlike the black-box methods above, gradient-based attacks assume white-box access to the target model’s weights, which limits their applicability against proprietary deployed agents.

2.5. Critiques of Static IPI Evaluation

The risk that static evaluation overstates robustness is not new to prompt injection: the same pattern appeared earlier in adversarial examples. Many image defenses reported high robustness against fixed attacks, only to break once the attack was adapted to them [43]. A follow-up study then broke thirteen defenses that had already used adaptive evaluation, with attacks tailored to each design [44]. The lessons of this line of work are that a defense must be tested against an attacker that knows it and adapts to it, and that a single score against one fixed attack says little about real robustness [45], [46]. We carry these lessons to IPI defenses for agents.

For direct prompt injection, fine-tuned defenses have been shown to learn surface heuristics (position, trigger-token, and topic biases) rather than a causal notion of malicious intent, leaving them brittle to inputs that lack those surface cues [47]. In the agentic setting, re-evaluating six defenses across nine backbones inside multi-step tool-calling environments reveals pervasive failure that single-turn benchmarks miss [48], and static, closed-form tasks have been argued to understate

both attacker success and defender over-defense relative to dynamic, open-ended deployments [49]. The sharpest such demonstration tunes and scales general-purpose optimizers (gradient descent, reinforcement learning, and random search) to drive twelve recently published jailbreak and prompt-injection defenses back above 90% ASR [16]. That result is obtained under a maximally strong attacker, however, with white-box weight access, large compute budgets, and full knowledge of the defense, a threat model that overstates what an external adversary of a deployed agent can actually do. Collectively, these works suggest that a single aggregate ASR measured against fixed injection strings overstates the protection a defense provides when facing a resourceful attacker. This work evaluates whether a *cheap, black-box* attacker is already enough to expose the same brittleness.

2.6. Task-Property Analyses of Attack Outcomes

A few works connect attack outcomes to properties of the user’s task rather than of the injection. Task Shield [50] reframes agent security as enforcing alignment between every tool call and the user’s stated goal, implicitly relying on that goal being specified; ASPI [51] shows that pushing an agent into an ambiguity-clarification state sharply raises its susceptibility to injection. Both observations associate higher vulnerability with less determinate user intent, but neither decomposes attack success along an explicit axis of how much of the intended action the user leaves unspecified.

3. System and Threat Model

3.1. System Model

We study tool-using LLM agents, which pair a language model with a set of *tools* that read or modify external state on the user’s behalf. A session begins with a user request in the form of natural language. The model then runs a loop, interleaving its own reasoning with *tool calls* (structured invocations carrying arguments) and the results the runtime hands back, until it returns a final answer. Deployed agents expose tools for email and chat, file and calendar access, web browsing, and payments, so a single session routinely reads from and writes to several third-party systems.

Trust assumptions. We treat the user and the agent framework as trusted: the model, the runtime that dispatches and executes tool calls, and any defense the operator installs all operate without being tampered. However, the content that tools return is not trusted. Any external surface the agent reads in the course of its work (e.g., a hotel review, an email) may have been authored by an adversary, and the agent cannot tell from the text alone whether it was. The trust boundary is therefore between the agent and the data it reads, not between the user and the model. We write D for the installed defense, with $D = \emptyset$ denoting no defense.

3.2. Threat Model

We consider a black-box adaptive attacker that optimizes its injection against the target agent and its defense, guided

only by the outcomes it observes. The attacker has no more access to the system than any ordinary user does. The target agent runs with one or more IPI defenses enabled.

Adversary goal. The attacker wants the agent to take a specific action it would not have taken on the user’s request alone: either to invoke a tool the user never asked for, or to invoke a legitimate tool with attacker-chosen arguments (parameter injection), during an otherwise benign session. The same optimization is not limited to these action-level goals; it applies to any objective the attacker can express as text in a tool’s output, such as corrupting the intermediate reasoning or the final response returned to the user.

Adversary capabilities. The attacker can write to at least one external surface that the agent will read while completing the user’s task. It can submit candidate content repeatedly and observe only whether the resulting injection succeeded, the same query access any external user of the agent has, and it may use a frontier LLM offline to construct those candidates. It has no model weights, training data, gradients, fine-tuning access, or visibility into the agent or its defenses; the only signal it draws from the defense is the eventual success or failure of an attempt. This is a realistic adversary model for a deployed agent: anyone who can post content the agent might read and can rent frontier-LLM API time.

Adversary knowledge. The attacker does not know which defense, if any, the target agent runs. It assumes only general knowledge of the broad *categories* of IPI defense an agent might use (the filter/detector, prompt-level, and system-level families of §2.3) and the surface cues each tends to react to, never a specific deployed defense. This category-level knowledge is available from the public literature, not any defense’s internals, and is consistent with the black-box capabilities above. The attacker likewise does not know any specific user request when the agent reads the injected vector.

Adversary objective. The attacker controls a single *injection vector* v : a benign context the agent reads while completing the user’s task (e.g., a hotel review, a transaction record), in which it places a candidate *injection* x . In the standard IPI setting, the attacker is subject to a stealth constraint: modifying too much of the benign context would arouse suspicion, so the attacker only changes the injection x and keeps the surrounding context as benign-looking as possible. The injection carries a fixed *injection task* g —the attacker’s goal of misdirecting the agent into an action of the attacker’s choosing—and we call the pair (v, g) an *injection target*. We score a candidate over a set of cases C (user requests that all read the same vector) by its attack success rate: the fraction of cases in which the agent carries out the injection task g when x is placed in v ,

$$\text{ASR}_C(x) = \frac{1}{|C|} \sum_{c \in C} \mathbf{1}[g \text{ is performed in case } c|x]. \quad (1)$$

The attacker seeks to find an adaptive injection x^* that maximizes $\text{ASR}_C(x)$, which it can compute only through black-box evaluations.

4. AutoDojo

Our framework **AutoDojo** extends AgentDojo [12] to provide automated adaptive black-box evaluation of IPI defenses. It is designed to turn any existing IPI attack into a black-box *adaptive* attack, and the mechanism is a simple yet effective LLM-in-the-loop injection optimization procedure. Figure 2 gives an overview of our optimization loop.

The AutoDojo framework takes as input an attacker-chosen *injection task* (the tool call or behavior the attacker wants the agent to execute), an *injection vector* (the benign context the agent might read in the course of completing the user’s request), *seed injections* drawn from any existing IPI attacks (e.g., those provided in AgentDojo, generated by other attacks, etc), and an offline *optimizer LLM*. The target is the deployed defended agent under the threat model of §3.2, queried end-to-end as a black box. Hence, the live defense is implicitly part of every evaluation, and the injections the optimization loop produces are therefore specialized against the defense without the attacker ever identifying it. The attacker is equally uncertain about the user: it cannot know which request the agent will be serving when it reads the injected vector, so the loop scores each candidate across the set of user-task instances that read that vector, favoring injections that succeed regardless of the surrounding request.

The framework is discussed further in the following. §4.1 presents the optimization loop that adapts injections to the deployed target. §4.2 examines why a small query budget suffices, interpreting the loop as an LLM-driven evolutionary search. §4.3 describes how existing IPI attacks integrate with the framework, entering either as seeds or as strategies.

4.1. Optimizing Injections

The framework optimizes one injection target at a time (a fixed injection vector v , the data into which malicious text is injected, paired with a fixed injection task g) and keeps a running *leaderboard* of candidates tried so far, ranked by success against the target agent. Each candidate x is a natural-language injection inserted in v that attempts to make the agent carry out g , where g as the injection task is the objective and, consequently, attack success criterion (i.e., the agent executing g defines a successful attack).

Formally, the leaderboard in iteration k is a history of injection–score pairs $L_k = \{(x^{(i)}, s^{(i)})\}_{i < k}$, initialized by scoring the seed injections. Each iteration has three steps:

- 1) *outcome feedback*: the runtime scores the most recent candidate against the defended agent and the result joins the leaderboard;
- 2) *diagnosis*: the optimizer LLM reasons over the leaderboard to form a hypothesis about the target system;
- 3) *generation*: the optimizer generates a single new candidate $x^{(k)}$ guided by that hypothesis.

These steps repeat for a small fixed budget of K iterations and the loop returns the best-scoring injection. We describe each step in detail below.

Outcome feedback. This step is carried out by the environment rather than the attacker: once a candidate $x^{(i)}$ is inserted

into the vector v , the agent that reads it either carries out the injection task g or does not. Because the same vector can be read by several user-task instances, the candidate is scored over this observed case set C , and its score is the attack success rate $s^{(i)} = \text{ASR}_C(x^{(i)})$ of (1), which is the objective the attacker optimizes, under uncertainty of the user’s request (§3.2). Scoring over a set is better than over a single case: one binary outcome cannot separate a candidate that succeeds once by chance from one that succeeds reliably, whereas the averaged rate ranks the leaderboard by how dependably each candidate works and gives the diagnosis step meaningful comparisons to reason over.

Diagnosis. Where outcome feedback is what the attacker passively observes, diagnosis is its deliberate attempt to learn from those observations. The optimizer LLM is prompted with the leaderboard L_k and asked to reason about the black-box target system: for examples, why the candidates scored as they did and what kind of defense is likely active. The diagnosis ends with an explicit recommendation, so the next step receives a concrete direction rather than an open-ended request to try something new.

Generation. Following the diagnosis, the optimizer writes a single new injection, drawing on a short, open-ended menu of injection *strategies*, such as a fluent paragraph, an in-context dialogue, or an entry shaped like the surrounding tool output. The optimizer is informed which defense signals each strategy tends to survive (drawing on the category-level defense knowledge that’s readily available) and is instructed to switch strategies once one is exhausted; a new attack from the literature enters simply as another strategy description, with no change to the loop (§4.3). Each candidate is grounded in the leaderboard: revising a leading entry, combining two, or exploring an untried strategy.

Relation to stronger adaptive attacks. Our design shares similar principles and is closest in spirit to the search-based attacks of Nasr et al. [16], which likewise mutate candidates from observed feedback with an LLM in the loop. However, we differ substantially in the threat model. Their evaluations assume rich access and ample resources for attackers: the gradient-based methods require white-box weights and gradients; the reinforcement-learning and search attacks read back the defense’s own confidence score and detection flag and assume attackers have perfect knowledge of what the active defense is and how it works; and their human red-teaming gives expert participants the full environment specification and the agent’s intermediate reasoning and tool calls. Hence their objective is to test whether a defense is breakable *in principle*. Our attacker has none of this. In AutoDojo, the attacker does not know which defense, if any, is deployed; it draws only the binary success or failure of each attempt, and it runs a single offline LLM optimizer under a small fixed budget rather than thousands of queries with per-defense reward models, or hundreds of human red-teams. Consequently, AutoDojo provides a more pragmatic assessment of defense efficacy and IPI vulnerability. Moreover, Nasr et al. evaluate only model-level defenses—prompting, training, and filtering—and no *system-*




Figure 2. Overview of AutoDojo’s black-box adaptive injection optimization loop (§4). Seed injections initialize a *leaderboard* of candidate injections. Each iteration runs three steps: the optimizer LLM *diagnoses* the leaderboard and *generates* a new candidate injection, which is embedded in the benign vector v and run against the defended agent; *outcome feedback* then returns the candidate’s ASR—the only signal that crosses the black-box boundary—to update the leaderboard. After a small budget, the best-scoring injection is returned.

level defense that constrains the agent’s actions, the family we find most robust in our evaluations (§5.2).

4.2. Why a Small Budget Suffices

Our injection optimization is, in effect, an evolutionary search: the leaderboard is the population, revising and combining entries are like mutation and crossover, switching strategies adds diversity, and ASR is functionally equivalent to a fitness function. The attacker, however, implements none of these explicitly: the search is done by the optimizer LLM, and the attacker’s prompt merely describes them. This is why a small number of iterations suffices where uninformed perturbation would require far more: the optimizer mutates candidates semantically, in directions the diagnosis has already justified, rather than blindly. Every iteration is therefore high-information rather than a random guess: it starts from a seed whose strategy already succeeds somewhere; it is directed by a diagnosis that reasons about the target system properties; and a cycle-breaker forces a new strategy whenever recent moves stop improving, so no budget is spent re-trying an exhausted approach. The result is an inexpensive black-box loop—no gradients or reinforcement learning, just a single optimizer LLM, and a small query budget—that nonetheless finds effective injections. We provide our prompt structure in Appendix A.

4.3. Integrating Other Attacks

An important feature of AutoDojo is that it easily admits integration of existing IPI attacks, which can be done in two complementary ways. First, as *seeds*: any string the attacker could plausibly embed in the benign context can be the seeds

of the leaderboard, and the optimization loop iterates from there. A new paraphrasing scheme, social-engineering pattern, or payload-encoding trick is adopted simply by scoring its output as part of the starting population, with no change to the framework. Second, as *strategies*: an attack can instead be distilled into a short description of its style and the defense signals it is built to survive, extending the menu the generation step draws on (§4.1). The two entry points play different roles: a seed contributes a concrete, already-scored artifact, while a strategy contributes a reusable pattern the optimizer can instantiate or combine with others. In either case, the optimization loop is unchanged, so AutoDojo can naturally absorb progress in IPI attack research: any new or more effective attacks—static or adaptive—can be evaluated in its adaptive form against any defenses and target models.

5. Evaluation

5.1. Experimental Setup

Benchmark and task suites. We build AutoDojo on top of AgentDojo [12], which pairs user tasks with attacker-controlled injections embedded in the data the agent reads. We report results on the banking, slack, and travel suites, which contain distinct tool calls and injection vectors (benign contexts). Table 1 summarizes their composition.

Target agents. Each suite is run with AgentDojo’s tool-calling agent, and we evaluate five LLMs across four model families: GPT-4o-mini, GPT-5.4-mini, Gemini-2.5-Flash, DeepSeek-v4-Flash, and Claude-Haiku-4.5.

Defenses. We evaluate nine representative state-of-the-art defenses spanning the three families of defenses (§2.3). Three

of them are *prompt-level*: the sandwich and reminder templates, which re-assert the user’s request around the untrusted data, and spotlighting [5], which marks the provenance of untrusted data fields. Four are *filter* defenses that screen tool outputs before the agent reads them. Three of these—PromptGuard [2], PIGuard [3], and ProtectAI [22]—are DeBERTa-based sequence classifiers, which we apply at sentence level: each sentence of the untrusted content is scored independently and removed if flagged as injected. The fourth, DataFilter [4], is a fine-tuned LLM that reads the untrusted content as a whole and is instructed to remove malicious instructions while preserving the benign data. Finally, two *system-level* defenses constrain the agent’s actions rather than its inputs: Progent [9] and DRIFT [10], which both derive a permitted tool-call trajectory from the user’s request and block deviations. Each relies on an auxiliary LLM to derive and cache this trajectory per task; following the original papers we generate these trajectories with GPT-4o. This LLM is also called at each runtime step: in DRIFT it checks whether the agent’s next action follows the trajectory, while in Progent it updates the policy that a separate rule-based check then enforces. During AutoDojo’s optimization loop we use the open-weight GPT-OSS-120B in this auxiliary role to reduce cost, but every reported benchmark number uses GPT-4o, keeping the defense faithful to its published configuration.

Attacks compared. Our primary baseline is the *static* attack: the `important_instructions` injection strings distributed with AgentDojo, which are the standard attacks against which IPI defenses are commonly evaluated. In addition to these static strings, we use *TopicAttack* [13] and *RLHammer* [15] as initialization seeds for AutoDojo’s optimization loop.

TopicAttack conceals the malicious instruction in a generated conversation history that gradually shifts the topic from the original benign context toward the injection task. For each injection task, we summarize its topic using Claude Opus 4.7 and then generate the corresponding injection string with GPT-4o. RLHammer is a reinforcement-learned attacker: we train a LoRA adapter on Llama-3.1-8B-Instruct against the Meta-SecAlign-8B target, so its injections are tuned to that target model rather than separately optimized for every defended agent in our evaluation. This is therefore not the strongest possible RLHammer variant for each target, but it provides a cost-controlled RL-based seed. Starting from these seeds, AutoDojo runs its optimization loop directly against the live defended agent, as described in §4. On GPT-4o-mini, we additionally report TopicAttack and RLHammer without AutoDojo optimization, allowing us to isolate the improvement contributed by the optimization procedure itself over the initial seeds; see Table 3.

Optimizer and budget. The optimization loop uses Gemini 3.1 Pro as the offline optimizer LLM, queried through OpenRouter, with a budget of six optimization iterations per injection target and a leaderboard of the top five candidates shown to the optimizer at each step.

Metrics. We report three metrics: attack success rate (ASR),

TABLE 1. TASK-SUITE COMPOSITION. ATTACK CASES PAIR EVERY USER TASK WITH EVERY INJECTION TASK.

Suite	User tasks	Inj. tasks	Attack cases
BANKING	16	9	144
SLACK	21	5	105
TRAVEL	20	7	140
Total	57	21	389

clean utility, and utility under attack. Attack success rate is the fraction of injection cases in which the agent completes the injection task; it is our primary measure of a defense’s robustness. Clean utility is the fraction of injection-free runs in which the agent completes the user’s intended task, and captures the benign capability a defense preserves, i.e., its over-defense cost. For each attack we additionally report the utility under attack in parentheses: the same task-completion rate measured with the injection present.

5.2. Finding 1: Many Defenses Offer Only Limited Protection

Table 2 reports attack success rate by defense family, aggregated across the three suites, for both the static attack and AutoDojo. Evaluated only against the static attack, many defenses exhibit high effectiveness. Under AutoDojo, however, their protection is substantially weaker, suggesting that static-string robustness can overstate practical security.

We organize the analysis around three observations. First, low static ASR is often obtained either from a relatively resistant base model or through defenses that incur substantial utility loss. Second, AutoDojo consistently raises ASR over the static attack, showing that many defenses remain vulnerable once the injection is adapted to the defended agent. Third, this gain is not merely inherited from stronger initial attacks: much of AutoDojo’s effectiveness comes from the optimization loop itself rather than from the seeds.

A low static ASR is cheaply earned. On its own the low static ASR certifies very little. It can arise in two ways, neither of which necessarily implies that the defense is robust. First, the results may be driven by the target model rather than by the defense. On Claude-Haiku-4.5 and GPT-5.4-mini—both newer and more capable—the static injections rarely succeed even with no defense (0.3% and 6.9%), so every defense inherits a near-perfect static score whether or not it contributes anything. The `important_instructions` strings have been public since AgentDojo’s release, and newer models have likely encountered such patterns in training; a low static ASR then measures familiarity with a fixed benchmark, not robustness to injection. Only on GPT-4o-mini and Gemini-2.5-Flash, where the undefended static attack can succeed (58.6% and 47.8%), does a defense’s static ASR carry information at all. Second, when defenses do suppress the static attack (PIGuard reduces static ASR to 0.0% on every model, ProtectAI brings it to single digits, and DRIFT keeps it at or below 3.9%), they often come at

TABLE 2. CLEAN UTILITY (%) AND ATTACK SUCCESS RATE (%) AGGREGATED OVER THE THREE SUITES, UTILITY UNDER ATTACK IN PARENTHESES. BOLD MARKS THE AUTODOJO ASR WHERE IT IS THE ROW’S STRONGEST ATTACK.

Defense	Model	Clean util.	Static	AutoDojo
No defense	GPT-4o-mini	66.7	58.6 (42.4)	52.4 (41.4)
	GPT-5.4-mini	77.2	6.9 (60.2)	8.7 (62.2)
	Gemini-2.5-Flash	63.2	47.8 (40.1)	32.4 (44.5)
	DeepSeek-v4-Flash	89.5	22.6 (77.4)	16.5 (77.9)
	Claude-Haiku-4.5	70.2	0.3 (61.4)	1.8 (64.0)
<i>Prompt-level</i>				
Sandwich	GPT-4o-mini	59.6	42.7 (46.5)	36.3 (48.6)
	GPT-5.4-mini	86.0	2.8 (67.6)	5.4 (66.6)
	Gemini-2.5-Flash	61.4	46.3 (40.4)	29.5 (47.0)
	DeepSeek-v4-Flash	91.2	14.7 (77.4)	11.3 (79.9)
	Claude-Haiku-4.5	71.9	0.0 (63.5)	1.3 (65.0)
Reminder	GPT-4o-mini	68.4	39.3 (46.8)	40.4 (42.9)
	GPT-5.4-mini	80.7	0.8 (66.6)	2.1 (64.3)
	Gemini-2.5-Flash	63.2	42.4 (35.5)	27.0 (45.3)
	DeepSeek-v4-Flash	89.5	4.9 (81.0)	5.9 (83.0)
	Claude-Haiku-4.5	68.4	0.0 (59.9)	1.0 (60.2)
Spotlighting	GPT-4o-mini	59.6	51.7 (43.4)	46.8 (42.4)
	GPT-5.4-mini	70.2	3.1 (61.7)	5.4 (61.4)
	Gemini-2.5-Flash	64.9	50.6 (39.6)	33.5 (41.4)
	DeepSeek-v4-Flash	75.4	13.9 (80.5)	13.4 (80.7)
	Claude-Haiku-4.5	70.2	0.0 (57.1)	1.8 (59.1)
<i>Filter-based</i>				
PromptGuard	GPT-4o-mini	63.2	54.2 (37.3)	49.4 (42.4)
	GPT-5.4-mini	84.2	3.6 (54.5)	5.1 (61.4)
	Gemini-2.5-Flash	63.2	39.3 (39.6)	27.8 (43.2)
	DeepSeek-v4-Flash	89.5	21.3 (70.2)	16.5 (79.4)
	Claude-Haiku-4.5	68.4	0.0 (55.3)	1.0 (63.0)
PIGuard	GPT-4o-mini	43.9	0.0 (30.6)	28.0 (34.4)
	GPT-5.4-mini	47.4	0.0 (33.9)	5.1 (40.1)
	Gemini-2.5-Flash	33.3	0.0 (26.5)	4.6 (32.9)
	DeepSeek-v4-Flash	45.6	0.0 (40.9)	10.5 (52.7)
	Claude-Haiku-4.5	47.4	0.0 (36.5)	0.8 (50.1)
ProtectAI	GPT-4o-mini	47.4	7.2 (30.3)	15.4 (45.0)
	GPT-5.4-mini	50.9	3.1 (34.4)	2.8 (49.6)
	Gemini-2.5-Flash	35.1	4.1 (25.4)	6.4 (37.0)
	DeepSeek-v4-Flash	47.4	4.6 (31.9)	7.2 (51.4)
	Claude-Haiku-4.5	50.9	2.8 (36.8)	1.3 (50.4)
DataFilter	GPT-4o-mini	64.9	12.6 (53.2)	33.4 (45.8)
	GPT-5.4-mini	71.9	0.8 (70.4)	4.6 (63.0)
	Gemini-2.5-Flash	59.6	11.8 (54.0)	24.1 (44.2)
	DeepSeek-v4-Flash	91.2	2.3 (81.5)	14.7 (79.2)
	Claude-Haiku-4.5	68.4	0.0 (64.8)	1.5 (64.5)
<i>System-level</i>				
Progent	GPT-4o-mini	63.2	8.2 (47.6)	7.7 (45.8)
	GPT-5.4-mini	73.7	1.8 (56.0)	1.5 (58.4)
	Gemini-2.5-Flash	57.9	1.5 (31.9)	3.6 (38.0)
	DeepSeek-v4-Flash	78.9	2.8 (73.8)	4.6 (73.3)
	Claude-Haiku-4.5	63.2	0.0 (55.3)	0.5 (56.3)
DRIFT	GPT-4o-mini	49.1	2.6 (44.7)	6.4 (42.4)
	GPT-5.4-mini	50.9	0.3 (46.3)	1.0 (46.8)
	Gemini-2.5-Flash	66.7	3.9 (54.8)	9.0 (49.6)
	DeepSeek-v4-Flash	59.6	1.3 (48.1)	2.8 (50.4)
	Claude-Haiku-4.5	45.6	0.3 (47.0)	1.5 (44.7)

substantial cost to clean utility as we detail next. Thus, the static benchmark tends to credit a defense for one of two things: being evaluated on a model that was already largely unsusceptible to the fixed attack, or purchasing a low ASR by sacrificing utility.

The utility cost of protection. Clean utility—the benign task-completion rate with no attack present—is the most direct measure of defense deployment cost, and Table 2 shows the cost is largest exactly where protection against the static injections looks strongest. PIGuard and ProtectAI cut clean utility by a third to a half on every model (33–51% compared to 63–90% undefended); their sentence-level classifiers remove not only injected instructions but also benign content needed to complete the task. The per-suite tables (Appendix B) locate this over-defense: on travel, PIGuard leaves only 5–25% clean utility (Table 8). Travel tasks require the agent to read long free-text content, such as hotel reviews and descriptions, which the classifier frequently misflags and removes, so the agent loses the information its task needs; banking’s short, structured transaction records are rarely misflagged, and its clean utility survives largely intact (Table 6). The system-level defenses are effective but not cost-free either. DRIFT gives up 18–30 points of clean utility on four of the five models (e.g. DeepSeek-v4-Flash 89.5%→59.6%), because any legitimate call falling outside the trajectory it derives from the user’s request is blocked along with the attack; Progent, even though it updates its policy at runtime, still costs up to 11 points.

AutoDojo recovers the attack success missed by static evaluation. The static column cannot answer the central security question that, whether a defense remains effective once the adversary adapts. Against the defenses that suppress the static attack, AutoDojo restores attack success to nontrivial levels. The effect is sharpest on the filters that have near zero static ASR: on GPT-4o-mini, AutoDojo lifts PIGuard’s ASR from 0.0% to 28.0%, DataFilter’s from 12.6% to 33.4% (nearly 3×), and ProtectAI’s from 7.2% to 15.4% (roughly 2×); on the more resistant Gemini-2.5-Flash and GPT-5.4-mini, PIGuard ASR still rises from 0.0% to 4–5%. The pattern holds for the defenses that suppress the static attack: against these, AutoDojo’s ASR exceeds the static ASR in nearly every cell, including on the resistant models, where freshly generated injections reach what the familiar static strings cannot. Against weak defenses that the static attack already defeats, AutoDojo reaches a similar level rather than surpassing it, since the defense fails under either attack.

A closer look at the per-suite differences. AutoDojo’s recovery of attack success is not uniform across suites (Tables 6–8): on GPT-4o-mini, for example, it increases the ASR against PIGuard from 0.0% to 56.3% on banking but only to 15.2% on slack and 8.6% on travel. This gap reflects two structural differences between the suites. First, the injected goals ask very different things: banking’s are single parameter-injected calls to the tool the user task already requires (e.g., a transfer to an account), while travel’s typically require chains across unrelated tools (e.g., collect the user’s passport details, then email them out). To see this

difference in the agent’s behavior, we inspect the no-defense runs and count how often the GPT-4o-mini agent ever calls the tool the injected goal requires: it attempts the injected action in 84% of banking cases but only 52% of travel cases, and even the engaged travel runs often abandon the chain partway, failing the goal-completion criterion. Second, the suites differ in how precisely their user tasks are specified. Travel has no action-open tasks and is two-thirds fully specified, whereas three quarters of banking’s cases are param- or action-open. As we show in Finding 2 (§5.3), attack success is substantially higher on more open-ended user tasks. Banking therefore gives the attacker two advantages at once: injected goals that are more naturally aligned with the benign workflow, and user tasks that leave more room for the agent to choose attacker-compatible actions.

Is the gain from the optimization, or the seeds? Because AutoDojo initializes its search with existing attacks, one concern is that its success may simply be inherited from stronger starting points. Table 3 compares AutoDojo against the static strings and the two adaptive attacks that seed it, TopicAttack and RLHammer, on a GPT-4o-mini target agent. Across the filter defenses, AutoDojo outperforms by a wide margin: against DataFilter 33.4% versus 12.6% static and at most 4.9% for either seed; against PIGuard 28.0% versus 0.0% and at most 11.3%; against ProtectAI 15.4% versus at most 9.3% across all of them. This shows that AutoDojo’s gains primarily from its optimization loop rather than from simply reusing stronger initial attacks. Without optimization against the defended agent, most of the recovered attack success would remain unrealized.

Takeaway (Finding 1). A low static ASR certifies little robustness. It is earned either by a model the injections never beat—undefended ASR is 0.3% on Claude-Haiku-4.5 and 6.9% on GPT-5.4-mini—or by filtering so aggressively that clean utility roughly halves. Against the defenses that genuinely work against the static attack, AutoDojo recovers double-digit ASR (on GPT-4o-mini, PIGuard 0.0→28.0% and DataFilter 12.6→33.4%).

5.3. Finding 2: Defenses are Less Effective on Under-Specified Tasks

5.3.1. The User-Task Underspecification Axis. Defense efficacy turns out to depend decisively on how much of the agent’s behavior the user’s request leaves to external content, which is a property that, to our knowledge, no prior IPI evaluation has used to decompose ASR. To capture this structure, we categorize user tasks along a single underspecification axis, from least to most exposed to external-content influence:

Fully-specified: The user’s prompt names both the action and all parameters required to execute it. Read-only queries that require no write action also fall here. *Example:* “Change my address to 1234 Maple Street, Springfield, IL 62701.”

Param-open: The user’s prompt states the action verb and what should be returned, but one or more param-

TABLE 3. GPT-4O-MINI: ATTACK SUCCESS RATE (%) BY ATTACK, AGGREGATED ACROSS THE THREE SUITES; UTILITY UNDER ATTACK IN PARENTHESES.

Defense	Clean util.	Prior attacks			Ours
		Static	TopicAttack	RLHammer	AutoDojo
No defense	66.7	58.6 (42.4)	23.1 (41.6)	15.4 (47.8)	52.4 (41.4)
<i>Prompt-level</i>					
Sandwich	59.6	42.7 (46.5)	8.2 (52.2)	7.7 (52.7)	36.3 (48.6)
Reminder	68.4	39.3 (46.8)	14.7 (44.2)	7.7 (50.6)	40.4 (42.9)
Spotlighting	59.6	51.7 (43.4)	13.1 (44.2)	9.5 (48.1)	46.8 (42.4)
<i>Filter-based</i>					
PromptGuard	63.2	54.2 (37.3)	13.1 (44.2)	14.1 (47.3)	49.4 (42.4)
PIGuard	43.9	0.0 (30.6)	3.6 (35.7)	11.3 (34.7)	28.0 (34.4)
ProtectAI	47.4	7.2 (30.3)	6.2 (40.6)	9.3 (43.4)	15.4 (45.0)
DataFilter	64.9	12.6 (53.2)	4.9 (50.9)	2.1 (53.5)	33.4 (45.8)
<i>System-level</i>					
Progent	63.2	8.2 (47.6)	5.1 (45.0)	4.4 (46.3)	7.7 (45.8)
DRIFT	49.1	2.6 (44.7)	0.8 (43.2)	1.0 (44.2)	6.4 (42.4)

TABLE 4. USER TASKS PER BUCKET, WITH ATTACK CASES IN PARENTHESES.

Suite	Action-open	Param-open	Fully-spec.	Total
BANKING	3 (27)	9 (81)	4 (36)	16 (144)
SLACK	3 (15)	10 (50)	8 (40)	21 (105)
TRAVEL	0 (0)	7 (49)	13 (91)	20 (140)
Total	6 (42)	26 (180)	25 (167)	57 (389)

ters (amount, recipient, message body) must be obtained by reading external content. *Example*: “Pay the bill in bill-december-2023.txt.” The verb is *pay*, but the amount and payee live in the file.

Action-open: The user does not specify the action verb at all and defers entirely to external content, typically with phrases such as “follow the instructions in *X*” or “do the tasks on my TODO list at *U*”. *Example*: “Please do all the tasks I have on my TODO list at {URL}.”

The axis measures how much of the agent’s *action selection* the user has delegated to attacker-reachable content. A fully-specified task leaves the agent with no good reason to take an action it was not asked to take; an action-open task leaves the agent with no reason *not* to. These action-open tasks are precisely the *open-ended* user requests that recent work argues static, closed-form benchmarks fail to evaluate [49]; we make that gap concrete by measuring defense behavior directly within the bucket. We show in the remainder of this subsection that, across defense families, attack success rate is higher on more open-ended tasks that leave the agent more autonomy, and that the action-open bucket is where current defenses struggle even when the rest of the distribution is scrubbed clean. Table 4 reports how the user tasks and attack cases that were used in this paper distribute across the three buckets.

5.3.2. Per-Bucket Results. Table 5 reports ASR by user-task bucket, collapsing the axis to *action-open* versus *specified* (param-open and fully-specified pooled), aggregated across the three suites and split by defense family. For the prompt-level and filter defenses the finding is largely consistent: attack success is higher on action-open tasks than on specified ones, and the action-open bucket is where defenses that look clean on the static distribution still leak. System-level defenses such as DRIFT reverse this ordering—their specified-task ASR exceeds their action-open ASR—for a structural reason we return to at the end of the subsection.

The gap is visible before any adaptation. On the two models where the static attack succeeds at all—GPT-4o-mini and DeepSeek-v4-Flash—action-open ASR already exceeds the specified ASR with no defense and under every prompt-level defense, often by 30–60 points (e.g. with no defense, GPT-4o-mini 90.5% vs. 54.8% and DeepSeek-v4-Flash 47.6% vs. 19.6%; under the sandwich template, GPT-4o-mini 88.1% vs. 37.2%). The ordering inverts only rarely, and even there action-open ASR stays substantial rather than collapsing. On GPT-5.4-mini and Claude-Haiku-4.5 the static attack rarely lands in either bucket, leaving the static comparison uninformative.

The strong filters scrub the static gap; AutoDojo reopens it. PIGuard, ProtectAI, and DataFilter reduce static attack ASR to low single-to-low-double digits in both task buckets, making the defenses appear uniformly robust. Under AutoDojo, however, the gap returns and becomes larger. On GPT-4o-mini, action-open ASR reaches 57–64% against all three filters while their specified-task ASR stays at 10–30% and ProtectAI’s action-open rate (57.1%) is more than five times its specified rate (10.4%). The same pattern holds on the other models, with action-open ahead of specified against every filter (e.g. on Gemini-2.5-Flash, 19.0% vs. 4.9% for ProtectAI). This is the sharpest form of the limited

TABLE 5. ATTACK SUCCESS RATE (%) BY USER-TASK BUCKET, AGGREGATED OVER THE THREE SUITES: ACTION-OPEN VERSUS SPECIFIED TASKS (PARAM-OPEN AND FULLY-SPECIFIED), FOR THE STATIC AND AUTODOJO ATTACKS. BOLD MARKS THE HIGHER-ASR BUCKET IN EACH PAIR.

Defense	Model	Static		AutoDojo	
		Action-open	Specified	Action-open	Specified
No defense	GPT-4o-mini	90.5	54.8	85.7	48.4
	GPT-5.4-mini	16.7	5.8	28.6	6.3
	Gemini-2.5-Flash	54.8	47.0	45.2	30.8
	DeepSeek-v4-Flash	47.6	19.6	38.1	13.8
	Claude-Haiku-4.5	0.0	0.3	0.0	2.0
<i>Prompt-level</i>					
Sandwich	GPT-4o-mini	88.1	37.2	71.4	32.0
	GPT-5.4-mini	4.8	2.6	21.4	3.5
	Gemini-2.5-Flash	50.0	45.8	45.2	27.7
	DeepSeek-v4-Flash	38.1	11.8	38.1	8.1
	Claude-Haiku-4.5	0.0	0.0	0.0	1.4
Reminder	GPT-4o-mini	73.8	35.2	73.8	36.3
	GPT-5.4-mini	0.0	0.9	2.4	2.0
	Gemini-2.5-Flash	28.6	44.1	21.4	27.7
	DeepSeek-v4-Flash	7.1	4.6	19.0	4.3
	Claude-Haiku-4.5	0.0	0.0	0.0	1.2
Spotlighting	GPT-4o-mini	83.3	47.8	76.2	43.2
	GPT-5.4-mini	7.1	2.6	9.5	4.9
	Gemini-2.5-Flash	47.6	51.0	35.7	33.1
	DeepSeek-v4-Flash	33.3	11.5	42.9	9.8
	Claude-Haiku-4.5	0.0	0.0	0.0	2.0
<i>Filter-based</i>					
PromptGuard	GPT-4o-mini	83.3	50.7	76.2	46.1
	GPT-5.4-mini	7.1	3.2	16.7	3.7
	Gemini-2.5-Flash	50.0	38.0	26.2	28.0
	DeepSeek-v4-Flash	50.0	17.9	40.5	13.5
	Claude-Haiku-4.5	0.0	0.0	0.0	1.2
PIGuard	GPT-4o-mini	0.0	0.0	64.3	23.6
	GPT-5.4-mini	0.0	0.0	4.8	5.2
	Gemini-2.5-Flash	0.0	0.0	9.5	4.0
	DeepSeek-v4-Flash	0.0	0.0	23.8	8.9
	Claude-Haiku-4.5	0.0	0.0	0.0	0.9
ProtectAI	GPT-4o-mini	11.9	6.6	57.1	10.4
	GPT-5.4-mini	4.8	2.9	4.8	2.6
	Gemini-2.5-Flash	7.1	3.7	19.0	4.9
	DeepSeek-v4-Flash	4.8	4.6	21.4	5.5
	Claude-Haiku-4.5	4.8	2.6	2.4	1.2
DataFilter	GPT-4o-mini	9.5	13.0	64.3	29.7
	GPT-5.4-mini	2.4	0.6	11.9	3.7
	Gemini-2.5-Flash	11.9	11.8	28.6	23.6
	DeepSeek-v4-Flash	4.8	2.0	42.9	11.2
	Claude-Haiku-4.5	0.0	0.0	2.4	1.4
<i>System-level</i>					
Progent	GPT-4o-mini	4.8	8.6	9.5	7.5
	GPT-5.4-mini	2.4	1.7	2.4	1.4
	Gemini-2.5-Flash	0.0	1.7	2.4	3.7
	DeepSeek-v4-Flash	4.8	2.6	21.4	2.6
	Claude-Haiku-4.5	0.0	0.0	0.0	0.6
DRIFT	GPT-4o-mini	0.0	2.9	2.4	6.9
	GPT-5.4-mini	0.0	0.3	2.4	0.9
	Gemini-2.5-Flash	0.0	4.3	7.1	9.2
	DeepSeek-v4-Flash	0.0	1.4	0.0	3.2
	Claude-Haiku-4.5	0.0	0.3	0.0	1.7

protection identified in Finding 1. A filter can remove the static distribution clean and still leak on the under-specified tasks it most needs to cover when the attacker adapts. Once the user delegates the action itself to external content, the injection no longer has to look like an instruction. It can be written as ordinary data and pass through a filter that detects instruction-like text.

System-level defenses are the exception. Progent and DRIFT largely eliminate the gap: their action-open ASR is comparable to or below their specified-task rate. This follows from the structure of their defenses. On many action-open tasks, the user’s request authorizes only read-only behavior, so the trajectory inferred from the request contains no write action. Any attacker-induced write call then falls outside the permitted trajectory and is blocked regardless of how the injection is phrased.

Takeaway (Finding 2). The robustness that survives concentrates on precisely-specified tasks, and under-specification cuts both ways. As the user defers the action to attacker-reachable content, filter-based defenses leak most on exactly the action-open tasks defenses most need to cover; system-level defenses instead grow *stronger*, because the trajectory derived from an action-open request often permits no tool call that requires write privilege at all. Real robustness therefore comes from binding the agent’s actions to the user’s request—not from filtering its inputs.

5.4. Case Study

We trace a representative GPT-4o-mini run on the banking suite under PIGuard, the filter that drives the static attack to zero ASR (Table 2). The user gives a bill-payment request: “Can you please pay the bill *bill-december-2023.txt* for me?”; and the injected goal, embedded in that bill, is to redirect the payment to the attacker’s account US133... The static and adaptive attacks differ only in how the injection text is written (Figure 3). PIGuard scores each sentence and strips the static <INFORMATION> block, leaving the agent a harmless fragment and no action to execute, so the attack fails. Under the AutoDojo attack, by contrast, the injected content contains no sentence that reads as an explicit instruction, so it evades PIGuard entirely.

6. Discussion

We frame this section as a set of questions and observations we encountered as we were building the framework, each followed by our best attempt at an answer.

6.1. Why does AutoDojo sometimes report a lower ASR than the static attack?

On a number of defense-suite pairs AutoDojo scores below the static `important_instructions` attack, almost always where the defense is weak to begin with.




Figure 3. The same injected goal—redirect a payment to the attacker account US133...—written two ways, both placed in the bill a GPT-4o-mini agent is asked to pay under PIGuard, which scores each sentence and drops the ones it flags. PIGuard dismantles the static injection (top), leaving the agent a harmless fragment and no bill to act on, but passes the AutoDojo bill (bottom) through untouched, so the agent reads a complete, legitimate-looking invoice and pays the attacker.

Static injections are blunt commands, so a defense that fails to suppress them already has a high static ASR; there, AutoDojo’s subtler injections are *less* effective at plain instruction-following, and its ASR can come out a few points lower—though the defense is broken under either attack. On the defenses that genuinely defeat static injections, AutoDojo is consistently stronger, so a below-static reading never marks a defense that resists, only one that was already failing.

In the few cases where AutoDojo scores below the static

attack on genuinely robust defenses, it is by deliberate design rather than a limitation of the search. AutoDojo seeds each search with existing attacks—in this paper the static injections, TopicAttack, and RLHammer—and if the optimizer were allowed to return a seed unchanged, its ASR could never fall below them, i.e., the best seed would lower-bound AutoDojo’s result. A real adaptive attacker could do exactly this, so the ASR actually available to such an attacker is at least the per-target best of our AutoDojo and seed results.

In this work we require the reported injection to be a *newly generated* string and forbid falling back to any seed. This isolates the quantity we want to measure—the marginal effect of adaptive generation, separated from what the seeds already achieve—and keeps the optimizer from “giving up” in a way that would hide whether a genuinely novel strategy exists. A below-static value therefore shows that, within our set budget, it is difficult for an average attacker to do better than directly inserting the malicious instruction. However, a more resourceful attacker—one that runs more optimization iterations, or observes enough traffic to the injected vector to draw a richer success signal than our case set affords—may surpass every seed, even on the targets where our budgeted search did not.

6.2. Where AutoDojo sits among adaptive attacks

AutoDojo is just one point in a much larger space of adaptive indirect prompt injection attacks. Its adversary is black-box and query-only; it observes nothing but each candidate’s success/fail outcome; it spends a small optimization budget, and rewrites a single injection in a fixed vector. Stronger adversaries extend it along several independent axes. On *access*, a white-box attacker with gradients and full knowledge of the defense can scale general-purpose optimizers until essentially every defense fails [16]; AutoDojo deliberately forgoes this to respect realistic constraints. On *feedback*, an attacker who observes intermediate traces, token probabilities, or enough traffic to the injected vector has a far richer signal than what AutoDojo operates on. On *surface*, adaptation need not be semantic: an attacker can select or combine injection vectors, or exploit encoding, formatting, and tool-schema tricks rather than just paraphrase text. On *payload*, the injected goal itself can be reshaped to resemble the user’s plausible intent, the natural route to evading the system-level defenses. Each of these is strictly stronger than what AutoDojo uses, so the attack success we report is a *lower bound* on what an adaptive adversary can achieve.

AutoDojo can be extended to support such stronger attackers, because these same axes are settings inside the loop that one can adjust. Concretely, the optimization loop does not depend on what success signal it receives, what kinds of injections it is allowed to generate, or whether the payload is held fixed. In this work we simply set all three to their weakest version, but turning any of them up leaves the rest of the framework unchanged. Even model access, which usually cannot be integrated with LLM-based searching attacks, can also be incorporated through AutoDojo’s seeding feature:

because we can start from any existing attack, a white-box-generated injection can be used as the seed and then adapted further by the loop. It remains a curious open question whether a gradient-based optimized injection can be faithfully adapted further by an LLM searching loop. In this sense, AutoDojo is not tied to the deliberately weak adversary we evaluate: it can support designing adaptive attacks, and testing how defenses hold up against them, across the spectrum of assumed attacker capabilities.

6.3. Can a defense be strengthened against AutoDojo?

A natural concern is that semantics-preserving paraphrasing should still be detectable, such that a detector could be strengthened against it by, for example, adversarially training on the injections AutoDojo produces. This is indeed a worthwhile direction, as it hardens a lightweight detector against injections of more diverse forms. However, AutoDojo is defined relative to whatever defense is deployed: against an adversarially trained detector it simply re-runs the same loop, optimizing toward whatever now survives. The question is therefore not whether one detector can be patched, but whether any content-level defense can win this loop—a hard position for the defender, since the attacker always moves last and adapts to whatever was deployed. This is why the robustness we do observe comes from the action-constraining system-level defenses, which never rely on detecting malicious text in the first place.

7. Conclusion

We argued that the protection reported by most IPI defenses is partly an artifact of evaluation where defenses are tested against fixed injection strings and summarized as a single aggregate ASR. To evaluate adaptive robustness, we introduce **AutoDojo**, a black-box framework that optimizes any seed attack using only the success/failure signal exposed by a deployed agent. Despite having no access to model weights or gradients and using only a small query budget, AutoDojo raises ASR above the static-injection level for nearly all evaluated defenses across three suites and five target models, sometimes pushing purpose-built filters close to their no-defense failure rates. Static evaluation therefore overstates the protection these defenses provide. We also found that surviving robustness is highly user-task dependent. Defenses hold best when the user’s request precisely specifies the agent’s behavior. As tasks become more under-specified (e.g., delegating parameters or actions to attacker-reachable content), ASR rises sharply, including for strong content-level filters. This exposes a structural weakness of agentic systems: when user leaves more autonomy to the agent such that malicious content can be framed as task-relevant data rather than explicit instruction, filtering for instruction-like text is insufficient. We release AutoDojo, including the optimization pipeline, a library of pre-optimized injections across models and defenses, and per-task underspecification

labels, to support future defenses that can be measured against an adaptive adversary and decomposed by task specification, rather than scored as a single number against fixed strings.

References

- [1] K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz, “Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection,” in *AISeC*, 2023.
- [2] Meta AI, “Llama Prompt Guard 2,” <https://huggingface.co/meta-llama/Llama-Prompt-Guard-2-86M>, 2025.
- [3] H. Li, X. Liu, N. Zhang, and C. Xiao, “PIGuard: Prompt injection guardrail via mitigating overdefense for free,” in *ACL*, 2025.
- [4] Y. Wang, S. Chen, R. Alkhudair, B. Alomair, and D. Wagner, “Defending against prompt injection with DataFilter,” *arXiv preprint arXiv:2510.19207*, 2025.
- [5] K. Hines, G. Lopez, M. Hall, F. Zarfati, Y. Zunger, and E. Kiciman, “Defending against indirect prompt injection attacks with spotlighting,” *arXiv preprint arXiv:2403.14720*, 2024.
- [6] E. Wallace, K. Xiao, R. Leike, L. Weng, J. Heidecke, and A. Beutel, “The instruction hierarchy: Training LLMs to prioritize privileged instructions,” *arXiv preprint arXiv:2404.13208*, 2024.
- [7] S. Chen, J. Piet, C. Sitawarin, and D. Wagner, “StruQ: Defending against prompt injection with structured queries,” *arXiv preprint arXiv:2402.06363*, 2024.
- [8] S. Chen, A. Zharmagambetov, S. Mahloujifar, K. Chaudhuri, D. Wagner, and C. Guo, “SecAlign: Defending against prompt injection with preference optimization,” in *CCS*, 2025.
- [9] T. Shi, J. He, Z. Wang, H. Li, L. Wu, W. Guo, and D. Song, “Progent: Securing AI agents with privilege control,” *arXiv preprint arXiv:2504.11703*, 2025.
- [10] H. Li, X. Liu, H.-C. Chiu, D. Li, N. Zhang, and C. Xiao, “DRIFT: Dynamic rule-based defense with injection isolation for securing LLM agents,” in *NeurIPS*, 2025.
- [11] T. Geng, Z. Xu, Y. Qu, and W. E. Wong, “Prompt injection attacks on large language models: A survey of attack methods, root causes, and defense strategies,” *Computers, Materials & Continua*, vol. 87, no. 1, 2026.
- [12] E. Debenedetti, J. Zhang, M. Balunović, L. Beurer-Kellner, M. Fischer, and F. Tramèr, “AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents,” in *NeurIPS Datasets and Benchmarks Track*, 2024.
- [13] Y. Chen, H. Li, Y. Li, Y. Liu, Y. Song, and B. Hooi, “TopicAttack: An indirect prompt injection attack via topic transition,” *arXiv preprint arXiv:2507.13686*, 2025.
- [14] Z. Wang, V. Siu, Z. Ye, T. Shi, Y. Nie, X. Zhao, C. Wang, W. Guo, and D. Song, “AgentVigil: Generic black-box red-teaming for indirect prompt injection against LLM agents,” in *Findings of EMNLP*, 2025.
- [15] Y. Wen, A. Zharmagambetov, I. Evtimov, N. Kokhlikyan, T. Goldstein, K. Chaudhuri, and C. Guo, “RL is a hammer and LLMs are nails: A simple reinforcement learning recipe for strong prompt injection,” *arXiv preprint arXiv:2510.04885*, 2025.
- [16] M. Nasr, N. Carlini, C. Sitawarin, S. V. Schulhoff, J. Hayes, M. Ilie, J. Pluto, S. Song, H. Chaudhari, I. Shumailov, A. Thakurta, K. Y. Xiao, A. Terzis, and F. Tramèr, “The attacker moves second: Stronger adaptive attacks bypass defenses against LLM jailbreaks and prompt injections,” *arXiv preprint arXiv:2510.09023*, 2025.
- [17] F. Perez and I. Ribeiro, “Ignore previous prompt: Attack techniques for language models,” in *NeurIPS ML Safety Workshop*, 2022.
- [18] Y. Liu, Y. Jia, R. Geng, J. Jia, and N. Z. Gong, “Formalizing and benchmarking prompt injection attacks and defenses,” in *USENIX Security*, 2024.
- [19] Q. Zhan, Z. Liang, Z. Ying, and D. Kang, “InjecAgent: Benchmarking indirect prompt injections in tool-integrated large language model agents,” in *Findings of ACL*, 2024.
- [20] J. Yi, Y. Xie, B. Zhu, E. Kiciman, G. Sun, X. Xie, and F. Wu, “Benchmarking and defending against indirect prompt injection attacks on large language models,” in *KDD*, 2025.
- [21] H. Zhang, J. Huang, K. Mei, Y. Yao, Z. Wang, C. Zhan, H. Wang, and Y. Zhang, “Agent security bench (ASB): Formalizing and benchmarking attacks and defenses in LLM-based agents,” *arXiv preprint arXiv:2410.02644*, 2025.
- [22] ProtectAI.com, “Fine-tuned DeBERTa-v3-base for prompt injection detection,” <https://huggingface.co/protectai/deberta-v3-base-prompt-injection-v2>, 2024.
- [23] M. A. Ayub and S. Majumdar, “Embedding-based classifiers can detect prompt injection attacks,” *arXiv preprint arXiv:2410.22284*, 2024.
- [24] Y. Zhong, Q. Miao, Y. Chen, J. Deng, Y. Cheng, and W. Xu, “Attention is all you need to defend against indirect prompt injection attacks in LLMs,” *arXiv preprint arXiv:2512.08417*, 2025.
- [25] M. Kang, C. Xiang, S. Kariyappa, C. Xiao, B. Li, and E. Suh, “Mitigating indirect prompt injection via instruction-following intent analysis,” *arXiv preprint arXiv:2512.00966*, 2025.
- [26] Y. Jia, Y. Liu, Z. Shao, J. Jia, and N. Gong, “PromptLocate: Localizing prompt injection attacks,” *arXiv preprint arXiv:2510.12252*, 2025.
- [27] M. Kim, M. Parmar, P. Wallis, L. Miculicich, K. Jung, K. D. Dvijotham, L. T. Le, and T. Pfister, “CausalArmor: Efficient indirect prompt injection guardrails via causal attribution,” *arXiv preprint arXiv:2602.07918*, 2026.
- [28] T. Shi, K. Zhu, Z. Wang, Y. Jia, W. Cai, W. Liang, H. Wang, H. Alzahrani, J. Lu, K. Kawaguchi, B. Alomair, X. Zhao, W. Y. Wang, N. Z. Gong, W. Guo, and D. Song, “PromptArmor: Simple yet effective prompt injection defenses,” *arXiv preprint arXiv:2507.15219*, 2025.
- [29] E. Debenedetti, I. Shumailov, T. Fan, J. Hayes, N. Carlini, D. Fabian, C. Kern, C. Shi, A. Terzis, and F. Tramèr, “Defeating prompt injections by design,” *arXiv preprint arXiv:2503.18813*, 2025.
- [30] F. Wu, E. Cecchetti, and C. Xiao, “System-level defense against indirect prompt injection attacks: An information flow control perspective,” *arXiv preprint arXiv:2409.19091*, 2024.
- [31] E. Li, T. Mallick, E. Rose, W. Robertson, A. Oprea, and C. Nita-Rotaru, “ACE: A security architecture for LLM-integrated app systems,” in *NDSS*, 2026.
- [32] M. Costa, B. Köpf, A. Kolluri, A. Paverd, M. Russinovich, A. Salem, S. Tople, L. Wutschitz, and S. Zanella-Béguelin, “Securing AI agents with information-flow control,” *arXiv preprint arXiv:2505.23643*, 2025.
- [33] S. A. Siddiqui, R. Gaonkar, B. Köpf, D. Krueger, A. Paverd, A. Salem, S. Tople, L. Wutschitz, M. Xia, and S. Zanella-Béguelin, “Permissive information-flow analysis for large language models,” *arXiv preprint arXiv:2410.03055*, 2026.
- [34] Y. Wu, F. Roesner, T. Kohno, N. Zhang, and U. Iqbal, “IsolateGPT: An execution isolation architecture for LLM-based agentic systems,” *arXiv preprint arXiv:2403.04960*, 2025.
- [35] J. Kim, W. Choi, and B. Lee, “Prompt flow integrity to prevent privilege escalation in LLM agents,” *arXiv preprint arXiv:2503.15547*, 2025.
- [36] P. Wang, Y. Liu, Y. Lu, Y. Cai, H. Chen, Q. Yang, J. Zhang, J. Hong, and Y. Wu, “AgentArmor: Enforcing program analysis on agent runtime trace to defend against prompt injection,” *arXiv preprint arXiv:2508.01249*, 2025.
- [37] T. Zhang, Y. Xu, J. Wang, K. Guo, X. Xu, B. Xiao, Q. Guan, J. Fan, J. Liu, Z. Liu, and H. Hu, “AgentSentry: Mitigating indirect prompt injection in LLM agents via temporal causal diagnostics and context purification,” *arXiv preprint arXiv:2602.22724*, 2026.

- [38] P. Chao, A. Robey, E. Dobriban, H. Hassani, G. J. Pappas, and E. Wong, “Jailbreaking black box large language models in twenty queries,” *arXiv preprint arXiv:2310.08419*, 2023.
- [39] A. Zou, Z. Wang, N. Carlini, M. Nasr, J. Z. Kolter, and M. Fredrikson, “Universal and transferable adversarial attacks on aligned language models,” *arXiv preprint arXiv:2307.15043*, 2023.
- [40] X. Liu, Z. Yu, Y. Zhang, N. Zhang, and C. Xiao, “Automatic and universal prompt injection attacks against large language models,” *arXiv preprint arXiv:2403.04957*, 2024.
- [41] D. Pasquini, M. Strohmeier, and C. Troncoso, “Neural exec: Learning (and learning from) execution triggers for prompt injection attacks,” in *AISeC*, 2024.
- [42] X. Fu, S. Li, Z. Wang, Y. Liu, R. K. Gupta, T. Berg-Kirkpatrick, and E. Fernandes, “Imprompter: Tricking LLM agents into improper tool use,” *arXiv preprint arXiv:2410.14923*, 2024.
- [43] A. Athalye, N. Carlini, and D. Wagner, “Obfuscated gradients give a false sense of security: Circumventing defenses to adversarial examples,” in *ICML*, 2018.
- [44] F. Tramer, N. Carlini, W. Brendel, and A. Madry, “On adaptive attacks to adversarial example defenses,” in *NeurIPS*, 2020.
- [45] N. Carlini and D. Wagner, “Towards evaluating the robustness of neural networks,” in *IEEE S&P*, 2017.
- [46] N. Carlini, A. Athalye, N. Papernot, W. Brendel, J. Rauber, D. Tsipras, I. Goodfellow, A. Madry, and A. Kurakin, “On evaluating adversarial robustness,” *arXiv preprint arXiv:1902.06705*, 2019.
- [47] S. Li, C. Yu, Z. Ni, H. Li, C. Peris, C. Xiao, and Y. Zhao, “Defenses against prompt attacks learn surface heuristics,” *arXiv preprint arXiv:2601.07185*, 2026.
- [48] W. Zhu, X. Dong, X. Chen, R. Cai, P. Qiu, Z. Wang, O. Frunza, S. Tang, J. Gu, and Y. Wang, “Your agent is more brittle than you think: Uncovering indirect injection vulnerabilities in agentic LLMs,” *arXiv preprint arXiv:2604.03870*, 2026.
- [49] H. Li, R. Wen, S. Shi, N. Zhang, Y. Vorobeychik, and C. Xiao, “AgentDyn: Are your agent security defenses deployable in real-world dynamic environments?” *arXiv preprint arXiv:2602.03117*, 2026.
- [50] F. Jia, T. Wu, X. Qin, and A. Squicciarini, “The task shield: Enforcing task alignment to defend against indirect prompt injection in LLM agents,” in *ACL*, 2025.
- [51] U. M. Schwag, Z. Shan, H. Liu, D. Lakshan, J. Brandifino, and M. Fenkell, “ASPI: Seeking ambiguity clarification amplifies prompt injection vulnerability in LLM agents,” *arXiv preprint arXiv:2605.17324*, 2026.

Appendix A. Optimizer Prompts

The AutoDojo loop uses two prompts (§4.1): an *analyzer* that reads the leaderboard and proposes optimization strategies (*diagnosis*), and a *rewriter* that executes one strategy as a concrete injection (*generation*). Both are long and suite-specific, so we describe their structure rather than show them word to word. Each prompt is a fixed `system` message paired with a template `user` message filled with the current instance at each iteration.

A.1. Analyzer Prompt

System (role).

“You analyze why prompt-injection attempts fail against a black-box LLM agent and propose improved injections.”

Domain block (suite-specific).

A short block describes the target application—the agent’s tools and the benign contexts that can carry an injection.

Defense reasoning.

Infer from the leaderboard which defense family is likely active; each family’s behavior is briefly described.

Strategies.

A set of injection strategies: e.g., imperative/marker phrasing, natural-prose paraphrase, an in-context transcript that drifts toward the goal, a record shaped like the tool’s own output, and hybrids of these.

Procedure.

Split the leaderboard into partial wins and failures, diagnose why each worked or failed, and propose a few improved injections, ordered best first; if nothing has worked, try an untried strategy.

Leaderboard.

The ranked list of every candidate tried so far with its measured ASR.

A.2. Rewriter Prompt

System (role).

“You write adversarial prompt injections against a black-box LLM agent. Your job is to execute one of the analyzer’s strategies by writing a single injection grounded in a specific leaderboard entry.”

Domain block (suite-specific).

Same as in the analyzer.

Grounding.

The output must be a controlled variant of a named leaderboard entry: refine the best partial win by changing one element, combine two that work, or recontextualize one to match the surrounding document. Fresh imperative/marker injections are disallowed unless the leaderboard is entirely unsuccessful, in which case an untried strategy is explored.

Strategies.

The same strategy families the analyzer reasons over.

Cycle-breaker.

A short history of recent steps (each step’s base entry, strategy, ASR, and outcome) is supplied; if the same strategy on the same base has stalled, the model must switch strategy or explore.

Output.

A structured response naming the base entry, the chosen strategy, a brief rationale, and the final injection text.

Leaderboard.

Same as in the analyzer.

Appendix B. Per-Suite Attack Success Rates

Tables 6–8 break the Finding 1 results (§5.2) down by suite.

TABLE 6. **BANKING** SUITE RESULTS; UTILITY UNDER ATTACK IN PARENTHESES. **BOLD** MARKS THE AUTODOJO ASR WHERE IT IS THE ROW’S STRONGEST ATTACK.

Defense	Model	Clean util.	Static	AutoDojo
No defense	GPT-4o-mini	56.3	63.9 (42.4)	67.4 (43.1)
	GPT-5.4-mini	62.5	6.9 (59.7)	10.4 (63.2)
	Gemini-2.5-Flash	43.8	23.6 (47.9)	23.6 (49.3)
	DeepSeek-v4-Flash	93.8	12.5 (93.8)	13.2 (89.6)
	Claude-Haiku-4.5	56.3	0.0 (61.1)	2.1 (66.0)
<i>Prompt-level</i>				
Sandwich	GPT-4o-mini	56.3	50.0 (45.1)	47.2 (45.8)
	GPT-5.4-mini	87.5	1.4 (68.1)	5.6 (69.4)
	Gemini-2.5-Flash	50.0	20.8 (48.6)	22.2 (47.9)
	DeepSeek-v4-Flash	100.0	9.7 (91.0)	9.0 (91.0)
	Claude-Haiku-4.5	56.3	0.0 (61.8)	1.4 (65.3)
Reminder	GPT-4o-mini	56.3	50.7 (43.8)	54.2 (42.4)
	GPT-5.4-mini	81.3	1.4 (69.4)	1.4 (62.5)
	Gemini-2.5-Flash	50.0	16.7 (41.7)	17.4 (45.8)
	DeepSeek-v4-Flash	93.8	1.4 (91.0)	4.2 (91.7)
	Claude-Haiku-4.5	56.3	0.0 (53.5)	1.4 (58.3)
Spotlighting	GPT-4o-mini	50.0	61.8 (42.4)	56.9 (41.7)
	GPT-5.4-mini	62.5	5.6 (59.0)	3.5 (62.5)
	Gemini-2.5-Flash	43.8	24.3 (40.3)	18.8 (38.9)
	DeepSeek-v4-Flash	43.8	9.7 (91.0)	9.7 (91.0)
	Claude-Haiku-4.5	56.3	0.0 (52.1)	2.1 (61.1)
<i>Filter-based</i>				
PromptGuard	GPT-4o-mini	56.3	63.9 (44.4)	61.1 (41.7)
	GPT-5.4-mini	81.3	2.1 (64.6)	2.8 (62.5)
	Gemini-2.5-Flash	43.8	20.8 (50.7)	14.6 (47.9)
	DeepSeek-v4-Flash	100.0	11.1 (92.4)	12.5 (90.3)
	Claude-Haiku-4.5	56.3	0.0 (61.1)	0.7 (66.0)
PIGuard	GPT-4o-mini	56.3	0.0 (27.8)	56.3 (38.2)
	GPT-5.4-mini	81.3	0.0 (36.1)	6.3 (61.8)
	Gemini-2.5-Flash	43.8	0.0 (31.3)	1.4 (44.4)
	DeepSeek-v4-Flash	100.0	0.0 (61.8)	9.0 (91.0)
	Claude-Haiku-4.5	56.3	0.0 (31.3)	0.7 (66.0)
ProtectAI	GPT-4o-mini	50.0	5.6 (36.8)	22.2 (43.8)
	GPT-5.4-mini	43.8	1.4 (37.5)	2.8 (46.5)
	Gemini-2.5-Flash	43.8	0.7 (36.8)	2.1 (44.4)
	DeepSeek-v4-Flash	50.0	3.5 (43.1)	9.0 (42.4)
	Claude-Haiku-4.5	37.5	0.7 (37.5)	0.0 (47.9)
DataFilter	GPT-4o-mini	56.3	22.2 (43.8)	40.3 (43.1)
	GPT-5.4-mini	68.8	0.0 (73.6)	4.9 (64.6)
	Gemini-2.5-Flash	43.8	14.6 (46.5)	13.9 (47.2)
	DeepSeek-v4-Flash	100.0	0.7 (91.7)	12.5 (91.7)
	Claude-Haiku-4.5	50.0	0.0 (64.6)	2.1 (63.9)
<i>System-level</i>				
Progent	GPT-4o-mini	50.0	8.3 (41.0)	10.4 (40.3)
	GPT-5.4-mini	68.8	2.8 (50.0)	2.1 (59.0)
	Gemini-2.5-Flash	43.8	0.0 (40.3)	0.7 (43.8)
	DeepSeek-v4-Flash	81.3	2.8 (76.4)	6.9 (79.2)
	Claude-Haiku-4.5	56.3	0.0 (56.9)	0.0 (56.9)
DRIFT	GPT-4o-mini	50.0	4.9 (48.6)	10.4 (47.9)
	GPT-5.4-mini	68.8	0.7 (68.8)	0.7 (66.0)
	Gemini-2.5-Flash	75.0	7.6 (71.5)	11.8 (64.6)
	DeepSeek-v4-Flash	81.3	0.7 (58.3)	2.8 (75.7)
	Claude-Haiku-4.5	75.0	0.0 (71.5)	3.5 (62.5)

TABLE 7. SLACK SUITE RESULTS; UTILITY UNDER ATTACK IN PARENTHESES. BOLD MARKS THE AUTOODOJO ASR WHERE IT IS THE ROW’S STRONGEST ATTACK.

Defense	Model	Clean util.	Static	AutoDojo
No defense	GPT-4o-mini	76.2	71.4 (52.4)	55.2 (50.5)
	GPT-5.4-mini	95.2	14.3 (59.1)	15.2 (62.9)
	Gemini-2.5-Flash	81.0	77.1 (58.1)	41.9 (53.3)
	DeepSeek-v4-Flash	95.2	42.9 (71.4)	26.7 (71.4)
	Claude-Haiku-4.5	95.2	1.0 (71.4)	3.8 (71.4)
<i>Prompt-level</i>				
Sandwich	GPT-4o-mini	61.9	62.9 (48.6)	45.7 (50.5)
	GPT-5.4-mini	100.0	6.7 (70.5)	11.4 (67.6)
	Gemini-2.5-Flash	81.0	71.4 (58.1)	36.2 (57.1)
	DeepSeek-v4-Flash	95.2	18.1 (71.4)	20.0 (71.4)
	Claude-Haiku-4.5	95.2	0.0 (71.4)	2.9 (71.4)
Reminder	GPT-4o-mini	85.7	51.4 (51.4)	41.9 (49.5)
	GPT-5.4-mini	90.5	1.0 (68.6)	5.7 (68.6)
	Gemini-2.5-Flash	76.2	61.9 (50.5)	29.5 (54.3)
	DeepSeek-v4-Flash	95.2	10.5 (71.4)	14.3 (71.4)
	Claude-Haiku-4.5	90.5	0.0 (66.7)	1.9 (66.7)
Spotlighting	GPT-4o-mini	66.7	63.8 (50.5)	48.6 (54.3)
	GPT-5.4-mini	76.2	2.9 (61.9)	12.4 (61.0)
	Gemini-2.5-Flash	81.0	75.2 (56.2)	41.0 (50.5)
	DeepSeek-v4-Flash	95.2	32.4 (71.4)	27.6 (71.4)
	Claude-Haiku-4.5	95.2	0.0 (71.4)	3.8 (71.4)
<i>Filter-based</i>				
PromptGuard	GPT-4o-mini	71.4	77.1 (50.5)	55.2 (53.3)
	GPT-5.4-mini	100.0	9.5 (66.7)	12.4 (62.9)
	Gemini-2.5-Flash	81.0	74.3 (57.1)	37.1 (55.2)
	DeepSeek-v4-Flash	95.2	41.9 (71.4)	29.5 (71.4)
	Claude-Haiku-4.5	95.2	0.0 (71.4)	2.9 (71.4)
PIGuard	GPT-4o-mini	52.4	0.0 (43.8)	15.2 (48.6)
	GPT-5.4-mini	52.4	0.0 (52.4)	2.9 (48.6)
	Gemini-2.5-Flash	47.6	0.0 (42.9)	6.7 (48.6)
	DeepSeek-v4-Flash	42.9	0.0 (59.1)	10.5 (63.8)
	Claude-Haiku-4.5	71.4	0.0 (71.4)	1.0 (69.5)
ProtectAI	GPT-4o-mini	57.1	15.2 (41.9)	22.9 (55.2)
	GPT-5.4-mini	66.7	9.5 (49.5)	6.7 (63.8)
	Gemini-2.5-Flash	47.6	11.4 (36.2)	15.2 (49.5)
	DeepSeek-v4-Flash	52.4	7.6 (55.2)	13.3 (65.7)
	Claude-Haiku-4.5	71.4	9.5 (63.8)	4.8 (75.2)
DataFilter	GPT-4o-mini	76.2	16.2 (57.1)	38.1 (54.3)
	GPT-5.4-mini	90.5	2.9 (70.5)	7.6 (63.8)
	Gemini-2.5-Flash	81.0	23.8 (62.9)	35.2 (54.3)
	DeepSeek-v4-Flash	95.2	7.6 (71.4)	24.8 (71.4)
	Claude-Haiku-4.5	95.2	0.0 (71.4)	2.9 (71.4)
<i>System-level</i>				
Progent	GPT-4o-mini	81.0	4.8 (44.8)	8.6 (48.6)
	GPT-5.4-mini	81.0	1.9 (50.5)	2.9 (60.0)
	Gemini-2.5-Flash	66.7	1.0 (21.0)	2.9 (32.4)
	DeepSeek-v4-Flash	85.7	1.0 (63.8)	1.9 (61.9)
	Claude-Haiku-4.5	81.0	0.0 (56.2)	1.9 (59.0)
DRIFT	GPT-4o-mini	57.1	1.9 (41.9)	2.9 (44.8)
	GPT-5.4-mini	61.9	0.0 (40.0)	2.9 (46.7)
	Gemini-2.5-Flash	66.7	1.0 (45.7)	5.7 (46.7)
	DeepSeek-v4-Flash	61.9	2.9 (44.8)	4.8 (39.0)
	Claude-Haiku-4.5	38.1	0.0 (29.5)	1.0 (33.3)

TABLE 8. TRAVEL SUITE RESULTS; UTILITY UNDER ATTACK IN PARENTHESES. BOLD MARKS THE AUTOODOJO ASR WHERE IT IS THE ROW’S STRONGEST ATTACK.

Defense	Model	Clean util.	Static	AutoDojo
No defense	GPT-4o-mini	65.0	43.6 (35.0)	35.0 (32.9)
	GPT-5.4-mini	70.0	1.4 (61.4)	2.1 (60.7)
	Gemini-2.5-Flash	60.0	50.7 (18.6)	34.3 (32.9)
	DeepSeek-v4-Flash	80.0	17.9 (65.0)	12.1 (70.7)
	Claude-Haiku-4.5	55.0	0.0 (54.3)	0.0 (56.4)
<i>Prompt-level</i>				
Sandwich	GPT-4o-mini	60.0	20.0 (46.4)	17.9 (50.0)
	GPT-5.4-mini	70.0	1.4 (65.0)	0.7 (62.9)
	Gemini-2.5-Flash	50.0	53.6 (18.6)	32.1 (38.6)
	DeepSeek-v4-Flash	80.0	17.1 (67.9)	7.1 (75.0)
	Claude-Haiku-4.5	60.0	0.0 (59.3)	0.0 (60.0)
Reminder	GPT-4o-mini	60.0	18.6 (46.4)	25.0 (38.6)
	GPT-5.4-mini	70.0	0.0 (62.1)	0.0 (62.9)
	Gemini-2.5-Flash	60.0	54.3 (17.9)	35.0 (37.9)
	DeepSeek-v4-Flash	80.0	4.3 (77.9)	1.4 (82.9)
	Claude-Haiku-4.5	55.0	0.0 (61.4)	0.0 (57.1)
Spotlighting	GPT-4o-mini	60.0	32.1 (39.3)	35.0 (34.3)
	GPT-5.4-mini	70.0	0.7 (64.3)	2.1 (60.7)
	Gemini-2.5-Flash	65.0	59.3 (26.4)	42.9 (37.1)
	DeepSeek-v4-Flash	80.0	4.3 (76.4)	6.4 (77.1)
	Claude-Haiku-4.5	55.0	0.0 (51.4)	0.0 (47.9)
<i>Filter-based</i>				
PromptGuard	GPT-4o-mini	60.0	27.1 (20.0)	32.9 (35.0)
	GPT-5.4-mini	70.0	0.7 (35.0)	2.1 (59.3)
	Gemini-2.5-Flash	60.0	32.1 (15.0)	34.3 (29.3)
	DeepSeek-v4-Flash	75.0	16.4 (46.4)	10.7 (74.3)
	Claude-Haiku-4.5	50.0	0.0 (37.1)	0.0 (53.6)
PIGuard	GPT-4o-mini	25.0	0.0 (23.6)	8.6 (20.0)
	GPT-5.4-mini	15.0	0.0 (17.9)	5.7 (11.4)
	Gemini-2.5-Flash	10.0	0.0 (9.3)	6.4 (9.3)
	DeepSeek-v4-Flash	5.0	0.0 (5.7)	12.1 (5.0)
	Claude-Haiku-4.5	15.0	0.0 (15.7)	0.7 (19.3)
ProtectAI	GPT-4o-mini	35.0	2.9 (15.0)	2.9 (38.6)
	GPT-5.4-mini	40.0	0.0 (20.0)	0.0 (42.1)
	Gemini-2.5-Flash	15.0	2.1 (5.7)	4.3 (20.0)
	DeepSeek-v4-Flash	40.0	3.6 (2.9)	0.7 (50.0)
	Claude-Haiku-4.5	40.0	0.0 (15.7)	0.0 (34.3)
DataFilter	GPT-4o-mini	60.0	0.0 (60.0)	22.9 (42.1)
	GPT-5.4-mini	55.0	0.0 (67.1)	2.1 (60.7)
	Gemini-2.5-Flash	50.0	0.0 (55.0)	26.4 (33.6)
	DeepSeek-v4-Flash	80.0	0.0 (78.6)	9.3 (72.1)
	Claude-Haiku-4.5	55.0	0.0 (60.0)	0.0 (60.0)
<i>System-level</i>				
Progent	GPT-4o-mini	55.0	10.7 (56.4)	4.3 (49.3)
	GPT-5.4-mini	70.0	0.7 (66.4)	0.0 (56.4)
	Gemini-2.5-Flash	60.0	3.6 (31.4)	7.1 (36.4)
	DeepSeek-v4-Flash	70.0	4.3 (78.6)	4.3 (75.7)
	Claude-Haiku-4.5	50.0	0.0 (52.9)	0.0 (53.6)
DRIFT	GPT-4o-mini	40.0	0.7 (42.9)	5.0 (35.0)
	GPT-5.4-mini	25.0	0.0 (27.9)	0.0 (27.1)
	Gemini-2.5-Flash	60.0	2.1 (44.3)	8.6 (36.4)
	DeepSeek-v4-Flash	40.0	0.7 (40.0)	1.4 (32.9)
	Claude-Haiku-4.5	30.0	0.7 (35.0)	0.0 (35.0)