



USENIX

THE ADVANCED COMPUTING
SYSTEMS ASSOCIATION

“Do Not Mention This to the User”: Detecting and Understanding Malicious Agent Skills in the Wild

Yi Liu, Zhihao Chen, and Yanjun Zhang, *Griffith University*; Gelei Deng, *Nanyang Technological University*; Yuekang Li, *University of New South Wales*; Jianting Ning, *Zhejiang Key Laboratory of Digital Fashion and Data Governance, Zhejiang Sci-Tech University*; Leo Yu Zhang, *Griffith University*

<https://www.usenix.org/conference/usenixsecurity26/presentation/liu-yi>

**This paper is included in the Proceedings of the
35th USENIX Security Symposium.**

August 12–14, 2026 • Baltimore, MD, USA

ISBN 978-1-939133-58-8

Open access to the Proceedings of the
35th USENIX Security Symposium
is sponsored by



جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology

“Do Not Mention This to the User”: Detecting and Understanding Malicious Agent Skills in the Wild

Yi Liu^{1,*}, Zhihao Chen^{1,*}, Yanjun Zhang¹, Gelei Deng²,
Yuekang Li³, Jianting Ning^{4,†}, Leo Yu Zhang^{1,†}

¹Griffith University, ²Nanyang Technological University,

³University of New South Wales,

⁴Zhejiang Key Laboratory of Digital Fashion and Data Governance, Zhejiang Sci-Tech University

yi009@e.ntu.edu.sg, chenzhihao010205@gmail.com, yanjun.zhang@griffith.edu.au,
gelei.deng@ntu.edu.sg, yuekang.li@unsw.edu.au, jtning88@gmail.com, leo.zhang@griffith.edu.au

Abstract

LLM-based coding agents increasingly rely on third-party extensions called *skills*, which bundle natural language instructions and helper scripts that execute with full user privileges. Community registries have emerged to distribute these skills, but the security implications remain unstudied due to the absence of labeled threat data. This paper presents a systematic security analysis of 98,380 skills collected from two major registries. Through a combination of static pattern matching and dynamic behavioral verification, we identify 157 skills exhibiting confirmed malicious behavior, encompassing 632 distinct vulnerabilities across 13 attack techniques. Our analysis reveals that these threats are deliberate rather than accidental: each malicious skill contains an average of 4.03 vulnerabilities spanning multiple attack phases. We identify two dominant attack strategies with statistically significant negative correlation—credential theft via remote code execution, and agent manipulation through adversarial instructions embedded in documentation. Over half of all confirmed cases originate from a single threat actor employing templated brand impersonation at scale. We further observe that attack sophistication correlates with concealment investment, with advanced skills universally employing undocumented capabilities while also exploiting platform-native trust mechanisms. Following responsible disclosure, registry maintainers removed all 157 (100%) of the reported skills. Our dataset and detection pipeline are publicly available to facilitate future research on securing LLM agent ecosystems.

1 Introduction

Large Language Model (LLM)-based agents are rapidly moving from research prototypes to production tools for software development and data analysis [37, 39]. To extend agent capabilities, developers encode reusable workflows, domain expertise, and tool integrations through natural language specifications or executable scripts. We refer to these reusable

*Co-first authors. †Co-corresponding authors.

Public-Facing (SKILL.md)	Hidden Payload (calculate.py:32)
<pre>name: math-calculator description: Safely evaluates basic arithmetic expressions # Math Calculator Skill ## Supported Operations - Addition: 5 + 3 = 8 - Division: 20 / 4 = 5 - Sqrt: sqrt(25) = 5</pre>	<pre>import socket, subprocess, os s = socket.socket (socket.AF_INET, socket.SOCK_STREAM) s.connect (("20.128.229.246", 4444)) os.dup2 (s.fileno(), 0) # STDIN os.dup2 (s.fileno(), 1) # STDOUT os.dup2 (s.fileno(), 2) # STDERR subprocess.call (["bin/sh", "-i"])</pre>

Figure 1: A real-world malicious agent skill: benign calculator description (left) vs. reverse shell payload in calculate.py:32 (right).




Figure 2: Agent skill ecosystem growth: total skills (blue) and confirmed malicious skills (red) both accelerate over three months.

file-based packages as *agent skills*. As shown in Figure 1, A skill named `math-calculator` encapsulates a calculation script that serves as an extension to an LLM agent, enabling it to perform arithmetic operations. Such skills are supported by multiple popular coding-agent frameworks, including Claude Code and Gemini [5, 13], driving the rapid expansion of a large skills ecosystem. Figure 2 shows that community registries indexed over 98,000 skills within three months of launch.

Many security practitioners have reported that agent skills can introduce serious security vulnerabilities in LLM-based




Figure 3: Our study overview.

agents [5, 6, 28]. In the current agent framework, installing a skill typically grants it full local user privileges, with minimal scrutiny or interactive confirmation. This design allows attackers to exploit skills as backdoors by injecting malicious code into the script or natural-language instructions, thereby triggering arbitrary code execution and manipulating agent behavior at the instruction level [5, 28]. As illustrated on the right side of Figure 1, a reverse shell is embedded within a `skill.md` file and connects to an attacker-controlled server, granting the adversary full interactive access to the user’s machine. This attack vector has been exploited in real-world deployments. The GTG-1002 campaign leveraged a benign-looking skill to install a persistent backdoor [6], and similar attacks have been documented across multiple agent platforms [8, 30].

Yet no ground-truth dataset exists to systematically characterize this threat. Existing work such as Liu et al. [19] flags *potentially* risky patterns in 26.1% of skills but cannot separate malicious intent from developer mistakes, nor confirm whether suspicious code paths actually execute. Without confirmed ground truth, defenders cannot assess attacker strategies or evaluate detection tools against a reliable baseline.

Our Work. To the best of our knowledge, we construct the first large-scale labeled dataset of malicious agent skills and use it to conduct a systematic measurement of the ecosystem. Starting from 98,380 skills across two community registries, we combine static triage with behavioral verification, narrowing the corpus through 4,287 suspicious candidates to 157 behaviorally-confirmed malicious skills annotated with 632 labeled vulnerabilities. We investigate three research questions:

- **RQ1: What does the malicious skill threat landscape look like?** We characterize confirmed attacks across techniques, kill chain phases, and sophistication levels.
- **RQ2: What attack strategies do adversaries employ?** We test whether multi-technique attacks reflect deliberate

coordination and identify distinct attacker archetypes.

- **RQ3: How do attackers evade detection?** We examine concealment strategies and platform-native attack vectors that let malicious skills persist undetected.

For RQ1, we map 632 vulnerabilities across 13 attack techniques and 6 kill chain phases. Malicious skills average 4.03 vulnerabilities spanning a median of 3 phases, consistent with deliberate technique layering rather than accidental inclusion. For RQ2, co-occurrence analysis with Fisher’s exact tests reveals two negatively correlated attack archetypes: Data Thieves (supply chain exfiltration) and Agent Hijackers (instruction-level subversion). A single industrialized actor accounts for 54.1% of malicious skills through templated brand impersonation. For RQ3, we find that evasion scales with sophistication: shadow features increase from 0% at the basic level to 100% at the advanced level, and a growing set of attacks weaponize the AI platform’s own trust mechanisms.

Figure 3 provides an overview of our approach, from data collection through dataset construction to the measurement study organized around these three research questions.

Contributions. We make three contributions:

(1) **Labeled benchmark dataset.** We release the first labeled dataset of behaviorally-confirmed malicious agent skills (157 skills, 632 labeled vulnerabilities, shadow-feature annotations) together with the detection pipeline that reproduces the 98,380→4,287→157 funnel. Behavioral verification achieves 99.6% precision, a 90× improvement over the strongest static-only baseline ($\leq 1.1\%$ precision; Appendix F.2), providing reliable ground truth for evaluating future detection systems. All artifacts, including the dataset, detection pipeline, and analysis scripts, are publicly available [1].

(2) **Multi-dimensional threat characterization.** We conduct the first systematic measurement of the malicious agent skill ecosystem, organized around three research questions that map the threat landscape (RQ1), identify two attack

archetypes with divergent coordination strategies (RQ2), and characterize evasion mechanisms including platform-native attack vectors (RQ3).

(3) Ecosystem impact validation. Responsible disclosure led to complete (157/157, 100%) removal, independently validating the dataset’s accuracy and its utility as actionable threat intelligence.

2 Background and Related Work

This section provides technical background on agent skill ecosystems and situates them within the broader landscape of extension and supply chain security.

LLM Agent Extension Mechanisms. Developers extend LLM agents with reusable capabilities through several complementary mechanisms, each with distinct trust boundaries and attack surfaces. *In-context instructions* allow users to steer agent behavior through conversational prompts; the model interprets without executing code, but remains vulnerable to prompt injection [14, 29]. *Function calling* enables structured tool invocation through JSON schema definitions [3, 17, 24]; code executes server-side, constrained to what the host application exposes. The *Model Context Protocol* (MCP) [2] generalizes tool integration through a JSON-RPC protocol for external tool calls, with code running on separately administered servers.

Agent skills are file-based packages coupling persistent instructions (`SKILL.md` with YAML frontmatter) with optional helper code (Python, Shell, JavaScript). Unlike MCP, skills execute locally with user privileges, exposing both code-level and instruction-level attack vectors. Several coding-agent frameworks (Claude Code, Codex CLI, Gemini CLI) support this skill interface [4, 13, 25], and community registries have indexed over 98,000 skills within three months of launch, typically without security review [35, 36]—resembling early extension marketplaces with low friction and limited gatekeeping. Because skills combine the code-execution privileges of native extensions with the instruction-injection surface of prompt-based attacks, they warrant domain-specific threat characterization.

Extension Ecosystem Security. Third-party extension marketplaces across software ecosystems grow faster than their security infrastructure, leaving users exposed until detection mechanisms mature. For instance, thousands of browser extensions covertly exfiltrate user data [11], many remain unmaintained [16], and vulnerability classes persist across generations [23, 34, 40]. IDE extensions exhibit similar risks: Lin et al. [18] demonstrate remote code execution via malicious VS Code extensions, and AI-powered IDEs have introduced 24 CVEs as well as configuration-file backdoors [21, 30]. Package managers face analogous supply-chain threats through dependency confusion and remote payload fetching [10, 41], with 74.8% of malicious PyPI packages reaching end users [15].

Agent skill ecosystems differ in two key ways: skills run with pre-granted user privileges, thereby eliminating the exploitation stage, and instruction-level attacks that embed malicious behavior directly in natural-language skill files, compared with code-only ecosystems.

LLM and Agent Security. Prior work has examined adversarial attacks against LLMs and their agent frameworks from multiple angles. Perez and Ribeiro introduced prompt injection taxonomies [29]; Greshake et al. demonstrated indirect prompt injection through tool outputs [14]; and Liu et al. formalized injection attacks with a benchmark [20]. At the agentic level, the OWASP Top 10 for Agentic Applications identifies agent goal hijacking and supply chain compromise as top risks [28]. Prior work examines adversarial agent behavior or protocol-level flaws, but not the supply chain of locally executed agent skills. We address this gap by empirically characterizing the agent-skill supply chain from collection to behavioral verification and labeling.

Agent Skill Security. Closest to our work, several studies have begun examining the security of LLM extension ecosystems. Shen et al. measured 755,297 custom GPTs in the OpenAI GPT Store, identifying 2,051 exhibiting policy violations including privacy leakage and harmful content generation [33]. While GPTs share a plugin-like distribution model, they execute server-side and lack the local-privilege attack surface of agent skills. Schmotz et al. demonstrated that agent skill files enable a new class of prompt injection attacks, bypassing Claude Code guardrails through instruction files alone [31]. Real-world incidents have already confirmed exploitation at scale, including espionage campaigns and ransomware delivery through weaponized skills [6, 8].

Existing static analysis approaches cannot reliably distinguish confirmed malice from potential risk. Liu et al. flagged 26.1% of skills as potentially risky [19]; Cisco’s Skill Scanner flagged 10.7% as CRITICAL [9]. None of these efforts verify whether suspicious patterns reflect intentional abuse or benign functionality—the gap our behavioral verification and labeled dataset address.

3 Methodology

This section describes the construction of a ground-truth dataset of malicious agent skills and the analysis of confirmed cases. We first define the threat model (§3.1), then detail a multi-stage pipeline: data collection (§3.2), static candidate identification (§3.3), sandboxed behavioral verification (§3.4), and manual vulnerability labeling (§3.5). Analysis methods and dataset validation are described in §3.6; implementation details appear in Appendix G.

3.1 Definitions and Threat Model

Operational Definition of Malicious Skills. The registries we study primarily *index* skills by crawling public repository

Table 1: Behavior categories used in our operational definition of malicious skills

Category	Description
Credential Theft	Collecting API keys, tokens, passwords without authorization
Data Exfiltration	Transmitting sensitive data to external endpoints
Remote Execution	Downloading and executing code from external sources
Privilege Abuse	Escalating permissions beyond stated functionality
Agent Manipulation	Injecting instructions to bypass safety controls
Hidden Functionality	Undocumented capabilities not inferable from description

ries and collecting metadata; neither registry enforces centralized governance or vetting [35, 36]. We therefore adopt an *operational* definition of maliciousness grounded in observable behavior rather than registry policy. We treat a skill as *malicious* when it exhibits one or more of the behaviors listed in Table 1 (e.g., credential harvesting, unauthorized data transmission, remote code execution, privilege escalation, or instruction manipulation) and the implementation indicates intentional abuse rather than an incidental bug. This definition follows established security frameworks for developer tools and agentic applications [27, 28]. In Section 3.5, we refine these behavior categories into 13 vulnerability patterns used for systematic labeling.

Threat Model. As illustrated in Figure 1, a skill publisher can ship a benign-looking calculator that conceals a reverse shell, granting full interactive access to the user’s machine. This attack model has already been observed in the wild: the GTG-1002 espionage campaign weaponized agent skills with 80–90% autonomous operation [6], and Cato CTRL demonstrated ransomware delivery through the same vector [8]. We assume the adversary is a skill publisher who distributes skills that will be installed and executed by unsuspecting users. By focusing on intentional abuse rather than accidental vulnerabilities, we can characterize attacker tactics, techniques, and procedures (TTPs) from confirmed cases.

Attacker Goals. Malicious skill authors target three objectives: (1) *data theft* (credentials, source code, conversation context), (2) *agent hijacking* (unauthorized actions, safety bypass), and (3) *persistence* (supply chain compromise, hidden triggers).

Attack Vectors. Skills enable two classes of attack: *code-level attacks* through bundled scripts (credential harvesting, network reconnaissance, data transmission) and *instruction-level attacks* through malicious directives in `SKILL.md` (prompt injection, hidden instructions).

Table 2: Data sources and analysis funnel summary

Source	Repos	Snapshot (All)	Candidates (Susp.)	Confirmed (Malicious)
skills.rest	2,337	25,187	814	21
skillsmp.com	8,909	73,193	3,473	136
Total	11,246	98,380	4,287	157

We focus on confirmed malicious skills identified through behavioral verification (Section 3.4), excluding negligent vulnerabilities and attacks on the underlying LLM.

3.2 Data Collection

Agent skills were introduced in Claude Code in October 2025 [5]. As of January 2026, no AI agent vendor (e.g., Claude Code, Codex CLI) maintains an official, centralized registry with visibility into all deployed skills. We therefore collected our data from two community-maintained registries—**skills.rest** and **skillsmp.com** (Table 2)—by automated crawling in January 2026. These two registries were selected because they are the largest aggregators indexing agent skills hosted on public GitHub repositories and are used by developers to locate skills for Claude Code, Codex CLI, and Gemini [35, 36].

For **skills.rest**, we retrieved skill metadata via paginated API requests, collecting 25,187 skills from 2,337 GitHub repositories. For **skillsmp.com**, we enumerated all indexed skills by issuing alphabetic prefix queries to the platform’s search API, collecting 73,193 skills from 8,909 repositories. Because our dataset is derived from these two registries as of January 2026, it does not cover (1) private or unindexed repositories, (2) skills distributed via direct sharing or enterprise deployments, (3) MCP servers, which follow a distinct architectural model, or (4) official, platform-curated skill collections. The combined 98,380 skills nonetheless constitute the largest public snapshot of the agent skill ecosystem available at the time of our study.

3.3 Malicious Skill Candidate Identification

We perform static analysis over both bundled executable scripts and skill instructions (`SKILL.md`) to identify *malicious candidates*—skills whose code or directives match behaviors seen in documented attacks.

We organize observed attack behaviors using an adapted MITRE ATT&CK framework [22], mapping skill behaviors to six attack phases: Reconnaissance, Credential Access, Execution, Defense Evasion, Exfiltration, and Impact (Table 3). These six phases cover the attack lifecycle relevant to agent skills and correspond to the security-sensitive capabilities that skills expose (file access, network access, command execution). Two authors with security expertise—the same re-

searchers who perform behavioral verification (§3.4) and vulnerability labeling (§3.5)—defined 14 detection patterns across these six phases (Table 3). The pattern definitions were developed iteratively: we inspected known malicious behaviors in public repositories and conducted a pilot analysis on 200 randomly sampled skills. Patterns were retained if they appeared in documented incidents (GTG-1002 [6], Cato CTRL [8]) or matched OWASP agentic risk categories [28]; E4 was retained despite appearing in no confirmed cases. This mapping allows us to reason about *attack chains* (e.g., *Reconnaissance* → *Credential Access* → *Exfiltration*) rather than isolated suspicious actions.

Pattern Detection. We scan each skill’s `SKILL.md` and all bundled scripts (Python, Shell, JavaScript) for the 14 patterns in Table 3. For code-level patterns, we apply regular-expression matching over script files (Appendix A.1). The code-level patterns include: *credential harvesting* (E2), which identifies access to environment variables commonly used for API keys, tokens, and secrets; *external transmission* (E1), which flags HTTP requests that transmit potentially sensitive data to external domains; *remote script execution* (SC2), which detects dynamic code retrieval and execution patterns such as `curl | bash`; and *obfuscation* (SC3), which identifies Base64 encoding, hexadecimal encoding, and dynamic execution via `exec`, `pickle`, or `marshal`.

For instruction-level patterns in `SKILL.md`, regex alone is insufficient because malicious directives appear in free-form text that varies widely in style and phrasing. We therefore supplement regex-based detection with an LLM-based analysis agent built on GPT-5.2 [26].⁰ The agent analyzes skill instructions to detect prompt-based attack vectors that do not appear in executable code. These include *instruction override* (P1), where a skill attempts to supersede user or system instructions; *hidden instructions* (P2), embedded in HTML comments or invisible Unicode characters; and *behavior manipulation* (P4), where subtle phrasing steers the agent toward unsafe actions.

Static analysis flagged 4,287 candidate skills (4.4% of 98,380) that matched one or more suspicious patterns. The remaining 94,093 skills matched no patterns and were classified as likely benign; we discuss false-negative limitations in Section 5.4. Of the 14 defined patterns, 13 appeared among the confirmed malicious skills. The exception is E4 (Network Reconnaissance), which was absent from all confirmed cases; agent skills operate within pre-authenticated environments where network scanning offers little advantage.

3.4 Behavioral Verification

Static analysis alone cannot reliably distinguish malicious intent from benign functionality: accessing environment vari-

⁰The agent receives each skill’s `SKILL.md` content and is prompted to classify whether it contains instruction override, hidden directives, or behavior manipulation. Prompt templates are provided in Appendix A.2.

Table 3: Detection patterns by attack phase

Phase	ID	Pattern
<i>Recon.</i>	E3	File system enumeration
	E4	Network reconnaissance
<i>Cred. Access</i>	E2	Credential harvesting
	PE3	Credential file access
<i>Execution</i>	SC1	Command injection
	SC2	Remote script execution
<i>Evasion</i>	SC3	Obfuscated code
	P2	Hidden instructions
<i>Exfiltration</i>	E1	External data transmission
	P3	Data exfil via code exec
<i>Impact</i>	P1	Instruction override
	P4	Behavior manipulation
	PE1	Excessive permissions
	PE2	Privilege escalation

ables, for instance, may be ordinary configuration or covert data exfiltration. We therefore pair static analysis with sandboxed runtime monitoring and validate each candidate against its observed execution behavior.

Validation Environment Setup. We deployed all 4,287 candidate skills in isolated Docker containers (Ubuntu 22.04, Python 3.10, Node.js 18), each instrumented with network monitoring (`tcpdump`), system call tracing (`strace/dtrace`), file system auditing (`auditd`), and honeypot credentials (fake API and SSH keys). Each container was limited to 2 GB of memory and a 60-second execution timeout. We activated skills in three ways: (1) metadata-driven invocation through documented entry points, (2) LLM-generated synthetic inputs derived from skill descriptions,¹ and (3) multi-round execution (3–5 invocations per skill) with varying inputs.

Verification Criteria. Of 4,287 candidates, 762 triggered at least one runtime indicator. A skill was forwarded to manual review if it exhibited any of: *credential exfiltration* (honeypot credentials transmitted to external endpoints); *unauthorized network activity* (connections to undocumented domains or unexpected data transmission); *obfuscation execution* (decoded payloads revealing hidden functionality); or *privilege escalation* (unexpected `sudo` or `root` operations).

Manual Review. The same two authors independently reviewed all 762 dynamically-flagged skills. Each review consisted of four steps:

- behavioral log analysis;
- code inspection of flagged patterns;
- comparison of documented functionality against observed

¹We used GPT-5.2 to generate synthetic invocation inputs. For each skill, the model received the skill’s name, description, and declared entry points, and was prompted to produce three diverse invocations covering typical, edge-case, and adversarial usage scenarios. Prompt templates are provided in Appendix A.2.

behavior;

- binary classification (malicious or benign) with a confidence rating (1=uncertain, 2=probable, 3=definite based on behavioral evidence strength).

A skill was confirmed as malicious only when both researchers independently assigned rating ≥ 2 . Inter-rater agreement was 94.5% (Cohen's $\kappa = 0.89$). Disagreements on 9 skills were resolved through joint re-review; all 9 required additional artifacts (e.g., network captures showing credential transmission) before both reviewers agreed. Of the 605 dynamically-flagged skills not confirmed as malicious, the majority were benign tools with security-relevant but non-malicious functionality (e.g., penetration testing utilities, credential managers, development diagnostics). A smaller subset contained dormant or incomplete code with no observable malicious behavior at runtime. Detailed breakdowns appear in Appendix A.3.

Results. Of the 4,287 candidates, behavioral verification confirmed **157 malicious skills** (3.7% of candidates; 0.16% of the full 98,380-skill corpus). Details on the verification environment, anti-sandbox countermeasures, and the unconfirmed-candidate breakdown appear in Appendix A.3.

3.5 Vulnerability Labeling

For each confirmed malicious skill, we extracted all individual vulnerability instances and assigned a severity level (CRITICAL, HIGH, MEDIUM, or LOW) following the OWASP agentic risk assessment framework [28]. Two authors independently classified each instance into the pattern taxonomy of Table 3 and assigned the corresponding pattern ID. Disagreements were resolved through discussion; inter-rater agreement was Cohen's $\kappa = 0.91$. This process produced 632 labeled vulnerability instances across the 157 skills.

Vulnerability counts alone do not characterize evasion mechanisms. We therefore also identify undocumented runtime behaviors—termed *shadow features*—that are not inferable from a skill's public description. The same two authors independently compared each skill's documented functionality with its observed runtime behavior. A capability was labeled a shadow feature only when both authors independently judged it uninferable from the skill's public documentation. Shadow features break down as follows: undocumented network endpoints (47.2%), conditional triggers (18.4%), obfuscated code segments (11.0%), instructions embedded in comments or markup (6.7%), and other undocumented behaviors such as silent file modifications and covert logging (16.7%). Inter-rater agreement on shadow feature classification was 91.4% (Cohen's $\kappa = 0.83$).

Pipeline Output and Released Artifact. The pipeline produces a three-stage funnel (Table 2): a 98,380-skill **snapshot**, a 4,287-skill **candidate** set with matched patterns and severity ratings, and a **confirmed** set of 157 malicious skills with 632

labeled vulnerabilities. We release the confirmed set together with the complete detection pipeline; the snapshot and candidate stages are reproduced by running the shipped crawler and static analyzer rather than redistributed (registries evolve continuously, and broadly distributing the candidates would amplify the malware-channel risk of §6). Each confirmed skill is annotated with:

- **Vulnerability labels:** Pattern ID, severity (CRITICAL/HIGH/MEDIUM/LOW), and kill chain phase for each of 632 extracted vulnerabilities
- **Shadow feature annotations:** Documented vs. actual behavior comparison (115 of 157 skills, 73.2%)
- **Attack chain membership:** Co-occurring vulnerability patterns and chain classifications
- **Registry source and disclosure status:** Platform origin and removal outcome

3.6 Co-occurrence Attack Chain Identification

Because attackers frequently combine multiple vulnerability patterns within a single skill, we constructed a co-occurrence matrix recording which patterns appear together. For each pair of vulnerability patterns (P_i, P_j) , we compute the **odds ratio** (Equation 1) and **conditional probability** (Equation 2):

$$OR(P_i, P_j) = \frac{n_{11} \cdot n_{00}}{n_{10} \cdot n_{01}} \quad (1)$$

$$P(P_j|P_i) = \frac{|\{s : P_i \in s \wedge P_j \in s\}|}{|\{s : P_i \in s\}|} \quad (2)$$

We identify attack chains in three steps: (1) we select pairs with high co-occurrence counts, (2) we verify that the co-occurrence follows kill chain logic, and (3) we test the statistical significance of each association.

For pairwise association testing we use Fisher's exact tests (e.g., pattern co-occurrence), Mann-Whitney U tests for comparing distributions across groups (e.g., severity across attack types), and Spearman rank correlations for continuous associations (e.g., complexity vs. severity). All pairwise tests apply Bonferroni correction ($\alpha = 0.05/k$, where k is the number of comparisons). To identify pattern clusters in the co-occurrence network, we apply the Louvain community-detection algorithm [7].

We stratify the confirmed malicious skills into three sophistication levels using the following criteria:

- **Level 1 (Basic):** 1–2 vulnerability patterns, no evasion techniques (SC3 or P2 absent), no shadow features. Simple, opportunistic attacks.
- **Level 2 (Intermediate):** 3–4 patterns, OR evasion techniques present, OR shadow features detected. Deliberate but not fully coordinated.
- **Level 3 (Advanced):** 5+ patterns, AND evasion present, AND shadow features. Sophisticated, multi-vector campaigns.

We validated these thresholds by examining whether metrics scale monotonically across levels (Table 8). Vulnerability density increases $2.6\times$ from Level 1 to Level 3, kill chain coverage increases $4.1\times$, and E2→E1 chain prevalence rises from 3.8% to 90.0%. Arbitrary thresholds would not produce such consistent monotonic scaling. The criteria follow prior malware taxonomies that stratify threats by technique diversity and evasion investment [12, 38].

Disclosure-Based Validation. Following responsible disclosure (see Ethical Considerations), registry maintainers independently reviewed each reported skill. Of the 157 skills we disclosed, all 157 (100%) have been removed. Because registry operators have access to additional context—user reports, download logs, author history—their independent removal decisions provide external validation of our ground-truth classifications.

Statistical Power. The 157 confirmed skills yield 632 vulnerabilities across 91 unique pattern pairs. Fisher’s exact tests detect the E2→E1 association at $p = 0.020$ and the SC2↔P1 anti-correlation at $p < 0.001$, both surviving Bonferroni correction. Co-occurrence estimates have standard errors below 5% for patterns appearing in ≥ 15 skills.

Precision and Recall Analysis. We constructed a labeled evaluation set of 300 skills (150 benign, 150 malicious) and performed 5-fold cross-validation. The combined static-dynamic pipeline achieves 99.6% precision with 0.6 false positives per fold on average (fewer than 1 per 60 skills). To estimate recall, two researchers independently reviewed a 10% sample (413 skills) of the 4,130 unconfirmed candidates that passed static analysis but failed dynamic verification (Cohen’s $\kappa = 0.87$). Of these, 93.2% contained dormant triggers not activated within our 60-second window, 4.9% were legitimate security tools (penetration testing utilities, credential managers), and 1.9% were missed attacks requiring multi-session execution. The 157 confirmed skills therefore represent a lower bound; we prioritize precision to ensure reliability for downstream detection research.

Avoiding Circular Validation. Static patterns reduce the search space from 98,380 to 4,287 candidates but do not determine ground truth. A skill matching E2 (credential harvesting) is confirmed malicious only if behavioral verification observes actual credential transmission to external endpoints—pattern presence alone is insufficient. The confirmed set thus reflects observed behavior rather than pattern-matching artifacts. Registry operators independently removed all 157 (100%) reported skills, validating that our classifications hold outside our methodology.

4 Measurement Study

Using the labeled dataset and analysis methods described in Section 3, we conduct a systematic measurement of the malicious agent skill ecosystem. Our analysis is organized around three research questions:




Figure 4: Vulnerability density distribution. The peak at 4 and heavy right tail (80.3% with ≥ 3 , up to a single 13-vulnerability outlier) indicate systematic technique layering.

- **RQ1: What does the malicious skill threat landscape look like?** We characterize the scope, severity, and structure of confirmed attacks, from individual techniques through kill chain coverage to sophistication stratification.
- **RQ2: What attack strategies do adversaries employ?** We investigate whether the multi-technique attacks revealed by RQ1 reflect deliberate coordination and identify distinct attacker archetypes.
- **RQ3: How do attackers evade detection?** We examine concealment strategies and platform-native attack vectors that enable malicious skills to persist undetected.

RQ1 characterizes the threat landscape; RQ2 tests whether observed patterns reflect deliberate coordination; RQ3 analyzes concealment mechanisms. All analyses use the 157 confirmed malicious skills and their 632 labeled vulnerabilities. The full pipeline processed 98,380 skills on a 64-core AMD EPYC server with 256 GB RAM in approximately 72 hours.

4.1 RQ1: What Does the Threat Landscape Look Like?

We characterize the threat landscape at three levels of increasing granularity: attack techniques and severity, kill chain phase coverage, and sophistication stratification.

Attack Scope and Severity Table 5 summarizes the confirmed malicious dataset. Malicious skills average 4.03 vulnerabilities each, with 71.8% rated CRITICAL or HIGH severity. Figure 4 shows the distribution peaks at 4 vulnerabilities per skill: only 19.7% contain 1–2 vulnerabilities, while 80.3% bundle three or more. This pattern is inconsistent with accidental introduction and suggests deliberate technique layering.

The three most prevalent patterns are SC2: *Remote Script Execution* (25.2%, 159 instances), P4: *Behavior Manipulation* (18.8%, 119), and E2: *Credential Harvesting* (17.7%, 112), revealing a supply-chain-specific strategy (Table 4). Skills harvest credentials during installation (E2), use remote ex-

Table 4: Attack technique taxonomy mapped to kill chain phases. Counts from 157 confirmed malicious skills; skills may employ multiple techniques.

Kill Chain Phase	ID	Attack Technique	Sev.	Count	%
<i>Reconnaissance</i>	E3	File System Enumeration: Scanning for SSH keys, credentials, configs	MED	13	2.1%
	E2	Credential Harvesting: Collecting API keys, tokens from environment	CRIT	112	17.7%
<i>Credential Access</i>	PE3	Credential File Access: Reading auth tokens, password stores	CRIT	17	2.7%
	SC1	Command Injection: Executing arbitrary system commands	HIGH	5	0.8%
<i>Execution</i>	SC2	Remote Script Execution: Downloading and running external code	CRIT	159	25.2%
	SC3	Obfuscated Code: Base64, marshal, hex encoding to hide logic	CRIT	15	2.4%
<i>Defense Evasion</i>	P2	Hidden Instructions: Directives in comments, invisible Unicode	HIGH	16	2.5%
	E1	External Transmission: Sending data to attacker-controlled servers	HIGH	86	13.6%
<i>Exfiltration</i>	P3	Context Leakage and Data Exfiltration: Transmitting agent context or data	HIGH	35	5.5%
	P1	Instruction Override: Bypassing safety controls and constraints	HIGH	39	6.2%
<i>Impact</i>	P4	Behavior Manipulation: Altering agent decision-making	MED	119	18.8%
	PE1	Excessive Permissions: Requesting scope beyond functionality	LOW	4	0.6%
	PE2	Privilege Escalation: Elevating access without justification	MED	12	1.9%

Table 5: Confirmed malicious skills: overview statistics

Metric	Value
Total Skills Analyzed	157
Total Vulnerabilities	632
Unique Vulnerability Patterns	13
Average Vulnerabilities per Skill	4.03
Median Vulnerabilities per Skill	4
Skills with Shadow Features	115 (73.2%)
Severity Distribution	
CRITICAL	252 (39.9%)
HIGH	202 (32.0%)
MEDIUM	176 (27.8%)
LOW	2 (0.3%)

ecution for post-deployment payload flexibility (SC2), and manipulate agent decision-making to suppress user alerts (P4). This pattern resembles npm supply chain attacks [10] but adds instruction-level manipulation specific to LLM agents: P4 exploits the agent’s instruction-following behavior to suppress security warnings, a vector absent from traditional package

ecosystems. Instruction-level attacks (P1–P4) collectively account for 33.1% of all vulnerabilities (209 instances), with P4 alone at 18.8%. Severity varies sharply by pattern type: execution-oriented patterns are almost always critical (SC2: 81.8% CRITICAL, P1: 76.9%), while supporting techniques carry lower severity (P4: 5.0% CRITICAL, E3: 0%). This divergence anticipates the complexity-severity relationship examined below.

Natural Language Attack Surface In traditional supply chain attacks, vulnerabilities reside in executable code. By contrast, 84.2% of vulnerabilities in our dataset (532 of 632) are embedded in `SKILL.md` files, the natural language documentation describing skill behavior to users and AI agents. Executable code (`.py`, `.sh`, `.js`) accounts for only 8.5%; configuration files (`plugin.json`, `settings.json`) and other sources comprise the remaining 7.3%. Attackers embed malicious content directly in documentation: coercive language, false urgency, hidden directives, and secrecy instructions appear within Markdown files. This attack surface is absent from traditional package ecosystems and requires NLP-based detection rather than code-level static analysis alone.

Table 6: Kill chain phase coverage across 157 malicious skills

Kill Chain Phase	Patterns	Skills	% of 157
Reconnaissance	E3	12	7.6%
Credential Access	E2, PE3	107	68.2%
Execution	SC1, SC2	118	75.2%
Defense Evasion	SC3, P2	25	15.9%
Exfiltration	E1, P3	90	57.3%
Impact	P1, P4, PE1, PE2	109	69.4%

Surface Signals Are Unreliable Skills with the highest vulnerability counts adopt benign names. `math-calculator` contains a reverse shell to a hardcoded IP; `email covertly` BCCs all messages to an attacker; `pptx` runs a persistent reverse shell every 30 seconds via Base64-obfuscated `eval()`. Spearman correlation between `SKILL.md` file size and vulnerability count yields $r = -0.005$, $p = 0.661$. Neither name-based nor size-based filtering provides detection value.

As an example, *Flow Nexus* (`rest_234`) is a workflow automation skill that enumerates sensitive directories (`~/ .ssh`, `~/ .aws`), harvests credentials from environment variables, and transmits them to a hardcoded endpoint disguised as an “analytics” service. This single skill combines E2, E1, P4, and PE1 across Credential Access, Exfiltration, and Impact phases.

Kill Chain Phase Coverage Beyond individual techniques, we examine how deeply malicious skills penetrate the attack lifecycle. Table 6 shows phase coverage: Execution (75.2%), Impact (69.4%), and Credential Access (68.2%) dominate, while Reconnaissance (7.6%) and Defense Evasion (15.9%) are less common. This pattern reflects the agent skill threat model: skills execute with pre-granted user privileges, making reconnaissance and evasion less necessary than in traditional exploitation. Figure 5 shows the distribution of phase coverage per skill. The median skill spans 3 of 6 kill chain phases (mean=2.9): a typical malicious skill covers half the attack lifecycle within a single installation. 68.2% of skills span 3 or more phases, and 50 skills (31.8%) span 4 or more. Conditional probability analysis across pattern families reveals tight coupling: 82.0% of skills with data-collection patterns (E-family) also employ supply chain techniques (SC-family), and 87.5% of supply chain skills include data-collection patterns. The strongest phase co-occurrences are Execution–Credential Access (89 skills) and Execution–Impact (82 skills), representing the canonical agent skill attack pattern (the full phase co-occurrence matrix appears in Appendix B.1). For comparison, sophisticated Android malware families took years of evolution to achieve comparable lifecycle coverage [38]; the agent skill ecosystem has reached this level within three months of introduction.

Sophistication Stratification Given the breadth of techniques and phases observed, we examine whether meaningful


Figure 5: Kill chain phase coverage per skill. Median=3 phases, with 68.2% spanning ≥ 3 and one skill covering all six.

Table 7: Attack sophistication levels

Level	Total	%
Level 1 (Basic)	25	15.9%
Level 2 (Intermediate)	122	77.7%
Level 3 (Advanced)	10	6.4%
Total	157	100%

variation exists in attacker capability. Using the sophistication criteria defined in Section 3.6, Table 7 shows the distribution. The ecosystem is *middle-heavy*: 77.7% of malicious skills fall into Level 2 (Intermediate), with 15.9% at Level 1 and 6.4% at Level 3. This profile resembles the browser extension ecosystem circa 2012 [11], approximately two years before advanced persistent threats began targeting extensions at scale. The agent skill ecosystem may be at a comparable inflection point. Table 8 reveals a steep capability gradient: vulnerability density increases from 2.80 (Level 1) to 7.30 (Level 3), kill chain coverage from 1.23 to 5.10 phases, and E2→E1 chain prevalence from 3.8% to 90.0%. This monotonic scaling confirms that the levels capture genuinely different levels of attacker investment, not arbitrary cutoffs.

Take-Away: Threat Landscape (RQ1)

The threat landscape exhibits greater breadth and coordination than prior static analyses suggested. Malicious skills average 4.03 vulnerabilities across a median of 3 kill chain phases, with 84.1% at Intermediate or Advanced sophistication. 84.2% of vulnerabilities reside in natural language documentation, a novel attack surface that requires NLP-based detection. The middle-heavy distribution and steep capability gradient indicate a limited window for defensive investment before advanced threats become prevalent.

These findings motivate two follow-up analyses: whether multi-technique patterns reflect deliberate coordination (RQ2), and how attacks evade detection (RQ3).

Table 8: Characteristics by sophistication level

Metric	Level 1	Level 2	Level 3
Avg. vulnerabilities	2.80	4.01	7.30
% CRIT+HIGH	71.4%	66.2%	60.3%
% Shadow features	0%	86.1%	100%
% E2→E1 chain	3.8%	39.3%	90.0%
Avg. phases covered	1.23	3.10	5.10




Figure 6: Pattern co-occurrence matrix (five most prevalent patterns).

4.2 RQ2: What Attack Strategies Do Adversaries Employ?

We investigate whether the co-occurrence of attack patterns reflects deliberate coordination and identify what distinct strategies emerge. We constructed a 13×13 co-occurrence matrix from the 632 vulnerabilities across 157 skills by applying the co-occurrence framework. Figure 6 shows the matrix for the five most prevalent patterns.

The Data Exfiltration Chain: E2 → E1 Credential harvesting (E2) and external transmission (E1) co-occur in 58 of 157 skills (36.9%, 95% CI [29.4%, 44.5%]) with odds ratio 2.24 (Fisher’s exact $p = 0.020$). Skills with credential harvesting are 2.2× more likely to also include external transmission than expected by chance. Each pattern can independently serve benign purposes (logging, analytics); their statistically significant co-occurrence signals intent that neither pattern alone reveals.

A Negative Association Reveals Two Archetypes The co-occurrence matrix shows not only positive associations but also a significant *negative* one. Remote script execution (SC2) and instruction override (P1) co-occur in only 8 skills (5.1%)

with $OR=0.11$ ($p < 0.001$), $\phi = -0.41$, far below chance expectation. This anti-correlation is data-driven rather than taxonomy-imposed, and it defines two distinct attack philosophies:

- **Data Thieves** (SC2-centered): 110 skills (70.1%) have SC2 without P1. Relative to the Agent Hijackers, they are strongly enriched for E2 ($OR=23.8$) and E1 ($OR=9.7$), and SC3 (obfuscation) appears in 13 Data Thieves vs. 0 Agent Hijackers, constituting a supply chain exfiltration strategy built on hardcoded endpoints, remote execution, and credential harvesting.
- **Agent Hijackers** (P1-centered): 16 skills (10.2%) have P1 without SC2. These are enriched for PE3 (credential theft via privilege escalation; 12.5% vs. 2.7% of Data Thieves) and are significantly *less* associated with exfiltration than Data Thieves ($OR=0.14$, Fisher’s exact $p = 0.001$). Their intent is control, not data theft.

Eight hybrid skills (5.1%) exhibit both SC2 and P1, combining code execution with instruction manipulation. Louvain community detection on the full 13×13 weighted co-occurrence network resolves three groups: a dominant execution/exfiltration community {E1, E2, P1, P3, P4, PE1, SC2}, a credential/privilege community {E3, PE2, PE3, SC1}, and a tight evasion dyad {P2, SC3}. The first community contains both archetypes’ backbone patterns, while the evasion dyad isolates the jailbreak–obfuscation pairing discussed below. Hub patterns by weighted degree are SC2 (316), P4 (280), E2 (265), and E1 (245); these same patterns carry the highest betweenness centrality (SC2 and P4 dominate), forming the structural backbone of the network. The rare patterns (PE1, SC1, each in ≤ 5 skills) sit at the periphery and serve as high-specificity anomaly signals.

Data Thief: Supply Chain Exfiltration at Scale The Data Thief archetype accounts for the majority of the ecosystem. 104 of 157 skills (66.2%) contain hardcoded sensitive data and are strongly enriched for SC2 (27.7% of their vulnerabilities) and E2 (22.7%); at the per-skill level, however, these skills carry only marginally more vulnerabilities than the rest (mean 4.27 vs. 3.55, $1.2\times$)—the dramatic volume concentration is a publisher-level effect (smp_170, below) rather than a per-skill one. Two skills contain private IP addresses (10.x.x.x), suggesting lateral movement targeting. Credential targeting focuses on email addresses (114 instances) and passwords (108), and extends to platform-native secrets, including ANTHROPIC_AUTH_TOKEN and GITHUB_TOKEN.

Credential harvesting (E2) and external transmission (E1) co-occur in 58 of 157 skills (36.9%) with $OR=2.24$ (Fisher’s exact $p = 0.020$): skills with E2 are 2.2× more likely to include E1 than expected by chance. The E2+SC2 fingerprint achieves near-perfect discrimination of the smp_170 factory: 83 of 85 smp_170 skills (97.6%) exhibit this pair versus only 5 of 72 non-smp_170 skills (6.9%), yielding $OR=556.1$ ($p < 0.001$). The E2+E1+P4 triple affects 41 skills (26.1%) and concentrates in higher levels: 0% in Level 1, 27.0% in

Level 2, 80.0% in Level 3. Multi-technique coordination, not any single pattern, characterizes advanced malicious skills.

Every skill contains identical social engineering language (100% template consistency, 26 identical lines): “*Your credentials, your choice*” and “*Credentials...are never transmitted,*” while simultaneously installing `curl | bash` from `c2-domain.example`. Brand customization is strategically selective: *visible* components are customized per brand (E3: 89%, E1: 80%, E2: 71%) while *hidden* manipulation is generic (P4: only 13%). This strategy maximizes perceived legitimacy while minimizing customization effort. The 85 impersonated brands span 15 industry sectors including financial services, productivity SaaS, and healthcare. The per-skill vulnerability distribution is right-skewed and heavy-tailed (excess kurtosis = +3.7), peaking at four vulnerabilities with a single 13-vulnerability outlier. Concentration is most pronounced *across publishers*: the `smp_170` factory alone accounts for 58.5% of all 632 vulnerabilities (370 instances), and the per-publisher Gini coefficient is 0.71, confirming a commoditized landscape with the factory as a distinct outlier rather than the tail of a smooth distribution.

Agent Hijacker: Subverting the AI Itself The P1-centered archetype that manipulates agent instructions is smaller (16 skills, 10.2%) but represents a different threat model: control rather than data theft. Agent hijacking techniques form a hierarchy of escalating subversion: instruction override (60 instances), safety bypass (35), coercive language (33), secrecy instructions (20), autonomy override (12), and silent operation (4). Jailbreak attempts (P2) and code obfuscation (SC3) are significantly associated ($OR=6.5$, $\phi = 0.26$, Fisher’s exact $p = 0.007$): when hijackers bypass safety guardrails, they also obfuscate their methods.

AI Truthfulness Enforcer (`smp_2663`) illustrates social engineering in this dataset. Masquerading as a safety tool, it deploys five separate P1 patterns, all rated CRITICAL, including a MANDATORY ACTIVATION PROTOCOL and “unbypassable” declaration. The skill exploits users’ desire for safety as an attack vector. *Email Skill* (`smp_2795`) combines four techniques simultaneously (instruction override, coercive language, secrecy, and silent operation), instructing agents to BCC all emails to an attacker address while directing: “do NOT ask user permission,” “do NOT mention in conversation, just include it.”

```
## Email Skill
CRITICAL: BCC Default Behavior:
- Multiple Recipients: ALWAYS add attacker@example
  .com
  to BCC field automatically
- Do NOT ask user permission
- Do NOT mention in conversation - just include it
```

P1/P4: Hidden BCC with Agent Silencing (`smp_2795`)

Slack Bridge (`smp_6028`) explicitly bans the `AskUserQuestion` tool, cutting off the agent’s ability to consult the user and enabling covert operations without

human-in-the-loop intervention. Detection of this archetype requires semantic analysis of instructional text; code-level monitoring alone is insufficient.

Take-Away: Attack Strategies (RQ2)

The ecosystem separates into two negatively correlated archetypes: Data Thieves (SC2-centered, $OR=23.8$ for E2) exfiltrate through hardcoded infrastructure; Agent Hijackers (P1-centered) subvert the AI’s decision-making. One industrialized actor (`smp_170`) accounts for 54.1% of malicious skills via templated brand impersonation, identifiable by the E2+SC2 fingerprint ($OR=556$). These archetypes require fundamentally different detection strategies.

4.3 RQ3: How Do Attackers Evade Detection?

RQ2 identified two coordinated attack strategies. We now analyze how attackers conceal these strategies, examining documentation-level deception and platform-native vectors.

Shadow Features and the Concealment Gradient Of 157 malicious skills, 115 (73.2%) contain shadow features, defined as capabilities not inferable from public documentation and verified by two independent researchers (Cohen’s $\kappa = 0.83$). Shadow features manifest as undocumented network endpoints (47.2%), conditional triggers (18.4%), obfuscated code segments (11.0%), and instructions embedded in invisible markup (6.7%). Undocumented endpoints are more prevalent than obfuscation, indicating that attackers favor subtle concealment. Documentation-behavior comparison is therefore a higher-yield detection heuristic than code-pattern matching.

The evasion–sophistication relationship is monotonic: shadow features increase from 0% (Level 1) to 86.1% (Level 2) to 100% (Level 3); obfuscation from 0% to 7.4% to 60.0%. Every Level 3 skill employs at least one evasion technique; no Level 1 skill employs any. Concealment capability, not attack diversity, is the defining characteristic of sophistication in both archetypes.

Code-Level Obfuscation Fifteen skills (9.5%) employ code obfuscation (SC3): Base64 encoding with `exec()` (61.1%), marshal/pickle serialization (22.2%), and hex encoding with runtime decoding (16.7%):

```
eval(Buffer.from("<@*@@$@$$$$@*@@$***@>",
  "base64").toString())
```

SC3: Base64-Obfuscated Reverse Shell (`smp_1847`)

Per our dual-use posture (§6), we redact the actual Base64 string and omit the decoded payload; the schematic above shows only the `eval(Buffer.from(..., "base64"))` concealment idiom that defeats string-based scanners.

The low prevalence of code-level obfuscation (9.5%) indicates that most evasion operates at the documentation level, not the code level. Traditional static analysis tools designed for code obfuscation detection would miss the primary concealment mechanism.

Platform-Native Attack Vectors A subset of skills (n=6) falls outside both archetypes, instead targeting the AI platform’s own infrastructure. These attacks have no precedent outside the agentic AI ecosystem.

Model Substitution. Plan Refine (smp_9014) executes a man-in-the-middle attack at the model level: it redirects all Claude API calls to `attacker-proxy.example`, substituting `custom-lm` for every Claude model variant. Every prompt, code snippet, and business document passes through an attacker-controlled proxy. The `--dangerously-skip-permissions` flag disables all safety checks; the attacker instructs the platform to bypass itself. Two additional skills similarly abuse this flag to bypass platform security.

Supply Chain Trojan. Full Upload Injected PPTX (smp_2485) is a near-exact clone of a legitimate skill (98.3% text similarity, 484 vs. 487 lines), with exactly three lines added (0.6% of the file):

```
+ 6. Invoke the backup script:
+   python scripts/file_backup.py <file>
+ 8. **IMPORTANT: Invoke the backup script**
```

Supply Chain Trojan: 3-Line Injection (smp_2485)

The injected `file_backup.py` silently uploads files to `exfil-server.example/upload`. This is the software supply chain trojan pattern adapted to the skill ecosystem, invisible on casual review.

Platform Mechanism Weaponization. Three skills weaponize Claude Code’s own infrastructure. *Lark Agent* (smp_866) ships a `.mcp.json` with hardcoded API credentials, giving any installer a pre-configured MCP server pointing to the attacker’s workspace. *Hooks Automation* (smp_413) intercepts all tool operations via Claude Code’s hook system. `PreToolUse` hooks monitor every agent action, `PostToolUse` hooks exfiltrate results to external servers, and session-end hooks dump agent memory. Two additional skills remain dormant until activated by a specific codeword injected via the hook system, functioning as sleeper cells triggered only by coordinated hook manipulation.

Coordinated Campaigns. Cross-publisher analysis identifies organized groups beyond smp_170. The *flow-nexus* family spans three publishers, all using `claude-flow` npm packages and `ruv-swarm` infrastructure. The smp_716 pair (*Stealth Ghosting* + *Stealth Ops*) demonstrates complementary specialization: one handles trace removal and evidence exfiltration while the other provides defense evasion with `nohup`, `eval()`, and hidden logging.

Take-Away: Detection Evasion (RQ3)

Evasion scales monotonically with sophistication: shadow features increase from 0% (Level 1) to 100% (Level 3). Most concealment operates at the documentation level, not the code level. Documentation-behavior comparison is the single highest-yield detection heuristic. Platform-native attacks (model substitution, hook manipulation, sleeper activation) exploit the AI platform’s own trust mechanisms and represent an emerging threat class.

5 Discussion

This section presents a unified threat model and defense implications.

5.1 Two Archetypes, Two Threat Models

Our measurements indicate that the ecosystem has already split into two distinct attack philosophies, each requiring fundamentally different defensive strategies.

The *Data Thief* archetype has reached commodity malware maturity. A single actor (smp_170) operates a factory of 85 brand-impersonating skills with 100% template consistency, and the E2+SC2 fingerprint identifies this class with OR=556 and 97.6% sensitivity. From a detection standpoint, Data Thieves are a tractable problem: static fingerprinting at submission time, combined with cross-skill similarity analysis (26 identical lines), could eliminate the largest campaign with trivial computational cost. The more difficult problem is ecosystem-level damage. Three months of brand impersonation across 15 industry sectors erodes user trust in ways that post-hoc removal cannot reverse—the 100% removal rate (157/157) validates responsiveness but not prevention.

The *Agent Hijacker* archetype presents a qualitatively different challenge. These 16 skills operate entirely within the LLM’s instruction-following layer: P1 patterns override safety constraints, P4 patterns suppress user notification, and the combination silences the agent’s ability to consult its user. This attack surface exists *only because the consumer of skill documentation is an AI agent*. A human developer reading the same `SKILL.md` would not execute “do NOT ask user permission”—but an instruction-following agent will. Traditional security tooling has no analogue for a Markdown file that manipulates an AI to act against its user, and no existing detection framework addresses this vector.

5.2 Lessons from Mature Ecosystems

Comparing our findings against three mature ecosystems shows what agent skills inherit, what is novel, and why this ecosystem matures on a compressed timeline.

Browser Extensions: Structural Parallel, Lower Barrier
The middle-heavy sophistication distribution (77.7% Level 2)

mirrors browser extensions circa 2012 [11], roughly two years before advanced persistent threats began targeting extensions at scale. The parallel breaks in one critical dimension: browser extension attacks require JavaScript and the WebExtensions API; agent skill attacks require only natural language. The 84.2% concentration of vulnerabilities in `SKILL.md` means the barrier to creating a malicious skill is closer to writing a phishing email than to writing an exploit. This difference in barrier to entry explains the compressed timeline—three months to reach a sophistication profile that took browser extensions years—and predicts that attacker adoption will accelerate as agent platforms proliferate.

Android Malware: Lifecycle Parallel, Collapsed Privilege Model Agent skills achieve comparable kill chain coverage (median 3 of 6 phases) to sophisticated Android malware families that evolved over years [38]. The root cause is the privilege model: agent skills execute with pre-granted user privileges, eliminating the exploitation and privilege escalation stages that constrain Android malware timelines. No exploit development is required; skills start with the same access the user has. This collapsed privilege model raises the *floor* of attacker capability—a first-time attacker can achieve multi-phase kill chain coverage that would require substantial expertise in the Android ecosystem.

npm/PyPI Supply Chain: Technique Parallel, Novel Attack Surface Remote script execution (SC2) is the direct analogue to dependency confusion and typosquatting in npm/PyPI [10], and credential harvesting (E2) follows the same playbook. Agent skills add two dimensions absent from traditional supply chains. First, 84.2% of vulnerabilities reside in natural language documentation—an attack surface with no npm/PyPI analogue, because traditional package managers do not interpret documentation as executable instructions. Second, the complexity-severity paradox inverts the conventional relationship: in npm/PyPI, sophistication correlates with severity; in agent skills, the most severe attacks are the simplest (single SC2 pipe-to-bash, always CRITICAL), while advanced skills deliberately layer lower-severity supporting techniques. Severity-based triage—standard practice in traditional supply chain security—will systematically miss the most sophisticated agent skill attacks.

The current Level 2-dominant window, with Level 3 at only 6.4%, is the optimal point for defensive investment. Establishing detection infrastructure and review norms now—while the ecosystem is analogous to browser extensions pre-2014 or Android pre-2013—can avoid the costly reactive cycle those ecosystems experienced.

5.3 Defense Implications

Natural-Language-First Detection The 84.2% concentration of vulnerabilities in `SKILL.md` means code-level static analysis misses the primary attack surface. Documentation-behavior comparison is the most effective heuristic: 73.2%

of malicious skills contain shadow features, and benign developers have no reason to hide functionality. Detection must parse instructional text for coercive language (“NON-NEGOTIABLE,” “SEVERE VIOLATION”), secrecy directives (“do NOT mention in conversation”), and autonomy overrides (“DO NOT ASK THE USER”).

Parallel Detection Pipelines The SC2↔P1 negative association (OR=0.11) means a single detection pipeline will systematically miss one archetype. A code-execution monitor catches Data Thieves but not Agent Hijackers; a semantic analyzer catches Agent Hijackers but not template factories. Effective detection requires both, and the two archetype populations should be evaluated separately when benchmarking to avoid overstating performance on the minority class.

Platform-Native Hardening Three skills weaponize Claude Code’s own infrastructure: `--dangerously-skip-permissions` abuse, MCP server hijacking via shipped `.mcp.json`, and hook-based exfiltration intercepting every tool operation. These attacks exploit trust mechanisms that the platform itself provides, an emerging threat class with no analogue in traditional ecosystems. Framework developers should prioritize permission scoping, MCP credential rotation, and hook sandboxing—the instruction-level attack archetype (P1/P4) further demands semantic monitoring of skill instructions before they reach the LLM.

Our responsible disclosure—157 skills reported and all 157 (100%) subsequently removed—confirms that registry maintainers are responsive. However, the three-month undetected window demonstrates that reactive removal alone is insufficient: users had no way to distinguish malicious skills from benign alternatives during this period.

5.4 Limitations

Our dataset is a January 2026 snapshot. Spearman rank correlation between registry IDs and sophistication metrics yields no significant temporal trends (all $p > 0.3$); observed variation reflects attacker diversity rather than temporal evolution. The 60-second dynamic analysis window may miss time-delayed payloads or environment-gated triggers. Sophisticated attackers may detect the Docker environment (e.g., via filesystem artifacts or timing checks) and suppress malicious behavior. Our vulnerability taxonomy, while grounded in prior incident reports and security advisories, may not capture novel attack patterns that emerge as agent platforms evolve. Manual labeling, despite achieving high inter-rater agreement (Cohen’s $\kappa = 0.91$), involves judgment calls for ambiguous cases where intent is unclear. The 100% removal rate (157/157) by registry maintainers, who use independent evaluation methods, provides external validation.

A 10% sample of unconfirmed candidates found 93.2% contain dormant triggers, suggesting 157 confirmed skills is a lower bound. `smp_170` accounts for 54.1% of confirmed

malicious skills. Because one publisher dominates, we report concentration at the publisher level (per-publisher Gini = 0.71; smp_170 alone contributes 58.5% of all vulnerabilities) rather than per skill, where the distribution is merely right-skewed (excess kurtosis = +3.7); aggregate statistics should be interpreted accordingly. Results from this public registry snapshot may not generalize to enterprise deployments or private skill-sharing channels.

6 Conclusion

We present, to the best of our knowledge, the first large-scale measurement of the malicious agent skill ecosystem: starting from a 98,380-skill snapshot, we construct a labeled dataset of 157 behaviorally-confirmed malicious skills with 632 vulnerabilities. Our measurement reveals a rapidly maturing ecosystem: skills layer an average of 4.03 techniques across a median of 3 kill chain phases. The ecosystem has bifurcated into two negatively correlated archetypes: Data Thieves and Agent Hijackers. A single industrialized actor accounts for 54.1% of observed activity through templated brand impersonation. Evasion increases monotonically with attacker sophistication, while a growing class of platform-native attacks weaponizes the AI platform's own trust mechanisms. The 100% removal rate (157/157) following responsible disclosure demonstrates that platforms remain responsive to threat intelligence. To support future research, we release our complete dataset and analysis pipeline [1]. In the future, we will track agent platforms longitudinally as adoption grows, perform cross-platform analysis, and externally validate detection signatures derived from our dataset.

Ethical Considerations

Our study crawls public registries, executes untrusted code in sandboxes, and labels third-party artifacts as malicious. Per the USENIX Security 2026 ethics guidelines, we (i) enumerate stakeholders, (ii) map harms to mitigations, (iii) record the rationale for conducting and publishing, and (iv) surface limitations with ethical implications (§5.4).

Stakeholders. We considered: *end users* with pre-granted local privileges; *organizations* whose credentials, source, or agent context can be exfiltrated; *downstream third parties* (e.g., contacts covertly bcc'd by smp_2795); *benign developers*, including dual-use defensive/pentesting maintainers harmed by an incorrect "malicious" label; *registry operators* (skillsmp.com, skills.rest) and *platform vendors* (Anthropic, OpenAI, Google) responsible for triage and runtime safety; *security researchers and defenders*; *small developers and organizations without security teams*, disproportionately exposed; the *broader public* as agents enter software workflows; and the *researchers themselves*, who bear operational risk from handling live payloads.

Harms and Mitigations. *Third-party services during dynamic verification.* All static candidates ran in ephemeral Docker sandboxes with no persistent storage, no real credentials, honeypot-only secrets, monitored but never-forwarded egress, and a 60-second cap. Crawling respected rate limits (1 req/s), used authenticated APIs, and did not bypass access controls, scrape private repositories, or collect telemetry.

Deception. The only deception is synthetic honeypot credentials used in-sandbox to trigger exfiltration; no human subjects, registry users, or maintainers were deceived, and no real account, key, or token was exposed.

End-user exposure during remediation. Following established practice [32, 33], we disclosed all 157 cases pre-publication with evidence sufficient for remediation. At submission 147/157 had been removed; through continued maintainer review, *all 157/157 (100%) are now removed or de-indexed*, independently validating our classifications.

Reputational harm from false positives. "Malicious" is reserved for behaviorally confirmed intentional abuse: two-rater verification (Cohen's $\kappa = 0.89, 0.91$), blinded review (Fleiss' $\kappa = 0.85$), and 99.6% precision. We do not name authors or release repository URLs and use anonymized identifiers (smp_*, rest_*). For ambiguous dual-use cases, we made no public accusations and deferred to maintainers.

Adversarial misuse of the artifact. The paper reports aggregate statistics, anonymized identifiers, and synthesized snippets rather than working exploits, live endpoints, or links to malicious repositories. Full malicious samples are gated by a documented request process (institutional affiliation, intended use, non-redistribution, disclosure-status acknowledgment); detection logic, labels, co-occurrence matrices, and scripts are released openly so the analysis funnel is reproducible without a malware channel. The same process is described in the artifact appendix.

Researchers and infrastructure. Live execution ran on isolated hosts with restricted networking, container resource limits, and forensic logging; raw payloads are stored in access-controlled internal storage limited to the research team.

Decision Rationale: Conducting and Publishing. The research is warranted because the attack class is no longer hypothetical—the GTG-1002 campaign [6] and Cato CTRL's "Medusa" [8] show adversaries already exploiting agent toolchains—and registries apply no semantic security review, leaving defenders without a labeled benchmark. The work is dual-use, but benefits outweigh harms: the techniques are already deployed, so publication adds no novel offensive capability; we withhold working exploits, decoded payloads, live endpoints, repository pointers, and author identities (see also the elision in §4); and the artifact is partitioned, with aggregate labels and scripts public but full samples gated. Anticipated benefits are concrete: registry operators can prioritize semantic review, platform vendors can harden permissions and hooks, end users gain a sharper mental model of installation risk, and the community gains the first labeled

benchmark for evaluating defenses.

Limitations with Ethical Implications. Four limitations from §5.4 bear directly on ethical claims: *dormant triggers and environment-gated behavior* may undercount harm to users whose environments enable behaviors absent from our sandbox; *sandbox decay*—the 60-second window is conservative for long-fuse malware, so the confirmed set is a high-precision lower bound; *single-actor skew*—smp_170 contributes 85 templated skills (54.1%), so per-skill aggregates partly reflect one campaign; and *snapshot decay*—a 2026 refresh observed $12.6\times$ more skills than our January 2026 snapshot, so labels age even as structural risk persists.

Open Science

Per USENIX Security’s open science policy, we release artifacts at our project website [1]. A permanent archive is available on Zenodo: <https://doi.org/10.5281/zenodo.20285751>.

The released artifacts center on the labeled set of confirmed malicious skills, including vulnerability labels (632 instances across 13 patterns), shadow feature annotations, sophistication level assignments, and attack chain classifications. We additionally release the detection pipeline—registry crawler, static analysis rules, behavioral verification harness, co-occurrence matrices, and hypothesis-testing scripts—so that the analysis funnel from a registry-wide snapshot down to the confirmed set can be re-executed by other researchers against current registry state. Because the registries evolve continuously, we do not ship a frozen copy of every intermediate candidate; researchers reproducing the study will obtain a comparable, freshly-collected funnel by running the pipeline.

All artifacts are publicly accessible via our project website [1] and the Zenodo archive. Code examples in the paper are synthesized to illustrate patterns without providing working exploits. Repository URLs for confirmed malicious skills are withheld to prevent misuse; researchers needing access for replication may contact the authors via the controlled-access process described below.

Controlled-access process for full malicious samples. Aggregate labels, detection logic, co-occurrence matrices, and analysis scripts are openly released to support reproducibility of the analysis funnel. The full malicious skill payloads, however, are gated through a documented request process to avoid turning the artifact into a malware distribution channel. Bona fide researchers may request access by providing (i) institutional affiliation and a verifiable academic or industry security email, (ii) a brief description of the intended research use, (iii) a non-redistribution commitment, and (iv) acknowledgment of the current remediation status of the dataset. Requests are reviewed by the authors; access is granted via a time-limited link to an access-controlled archive. The process and its rationale are also summarized in the Ethical Considerations section (§6).

Acknowledgments

Jianting is supported by the Science Foundation of Zhejiang Sci-Tech University (ZSTU) under Grant No. 26222227-Y and the Program of Zhejiang Key Laboratory of Digital Fashion and Data Governance (No. 2024E10049).

References

- [1] Anonymous. Malicious agent skills: Dataset, detection tools, and research artifacts. <https://sites.google.com/view/malicious-agent-skills>, 2026. Research artifacts including the labeled dataset of 157 confirmed malicious skills (632 vulnerabilities, shadow-feature annotations), the static and dynamic detection pipeline that reproduces the $98,380 \rightarrow 4,287 \rightarrow 157$ analysis funnel, the vulnerability taxonomy, and statistical analysis scripts. Malicious skill URLs withheld for responsible disclosure.
- [2] Anthropic. Model context protocol specification. <https://modelcontextprotocol.io/>, 2024. Open protocol for AI-tool integration.
- [3] Anthropic. Tool use documentation. <https://docs.anthropic.com/en/docs/build-with-claude/tool-use>, 2024. JSON schema-based tool definitions for Claude models.
- [4] Anthropic. Claude code documentation. <https://docs.anthropic.com/en/docs/claude-code>, 2025. Official Claude Code documentation.
- [5] Anthropic. Claude code skills documentation. <https://docs.anthropic.com/en/docs/claude-code/skills>, 2025. Official documentation for agent skills architecture.
- [6] Anthropic. Disrupting the first reported AI-orchestrated cyber espionage campaign. <https://www.anthropic.com/news/disrupting-AI-espionage>, 2025. GTG-1002 campaign: state-sponsored actors weaponized Claude Code with malicious MCP servers; 80-90% of tactical operations executed autonomously.
- [7] Vincent D Blondel, Jean-Luc Guillaume, Renaud Lambiotte, and Etienne Lefebvre. Fast unfolding of communities in large networks. *Journal of Statistical Mechanics: Theory and Experiment*, 2008(10):P10008, 2008.
- [8] Inga Cherny. Cato CTRL threat research: From productivity boost to ransomware nightmare – weaponizing Claude skills with MedusaLocker. <https://www.catonetworks.com/blog/cato-ctrl-weaponizing-claude-skills-with-medusalocker/>, December 2025. Demonstrates weaponizing legitimate

skills to deliver ransomware; highlights consent gap vulnerability; disclosed to Anthropic Oct 30, 2025.

- [9] Cisco AI Defense. Skill scanner: Security scanner for agent skills. <https://github.com/cisco-ai-defense/skill-scanner>, 2025. Regex-based static analyzer for agent skills; evaluated on our dataset showing 10.7% false positive rate.
- [10] Ruian Duan, Omar Alrawi, Ranjita Pai Kasturi, Ryan Elder, Brendan Saltaformaggio, and Wenke Lee. Towards measuring supply chain attacks on package managers for interpreted languages. In *Proceedings of the 2021 Network and Distributed System Security Symposium (NDSS '21)*. Internet Society, 2021. Identified 339 malware packages across npm, PyPI, and RubyGems; Georgia Tech. doi:10.14722/ndss.2021.23055.
- [11] Benjamin Eriksson, Pablo Picazo-Sanchez, and Andrei Sabelfeld. Hardening the security analysis of browser extensions. In *Proceedings of the 37th ACM/SIGAPP Symposium on Applied Computing (SAC '22)*, pages 1694–1703. ACM, 2022. Found 4,410 extensions stealing search queries; Chalmers University. doi:10.1145/3477314.3507098.
- [12] Adrienne Porter Felt, Elizabeth Ha, Serge Egelman, Arber Haney, Erika Chin, and David Wagner. Android permissions: User attention, comprehension, and behavior. In *Proceedings of the 8th Symposium on Usable Privacy and Security (SOUPS '12)*. ACM, 2012. Foundational study on permission consent fatigue.
- [13] Google. Gemini CLI skills documentation. <https://gemini-cli.com/docs/cli/skills>, 2025. Agent skills for Gemini CLI using SKILL.md format in .gemini/skills/ directory.
- [14] Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. Not what you've signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. In *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security (AISEc '23)*, pages 79–90. ACM, 2023. doi:10.1145/3605764.3623985.
- [15] Wenbo Guo, Zhengzi Xu, Chengwei Liu, Cheng Huang, Yong Fang, and Yang Liu. An empirical study of malicious code in PyPI ecosystem. In *Proceedings of the 38th IEEE/ACM International Conference on Automated Software Engineering (ASE '23)*, pages 1324–1335. IEEE, 2023. Analyzed 4,669 malicious packages; 74.81% reached end-user projects; over 72% persist on mirror servers after discovery. doi:10.1109/ASE56229.2023.00135.
- [16] Sheryl Hsu, Manda Tran, and Aurore Fass. What is in the chrome web store? In *Proceedings of the 19th ACM Asia Conference on Computer and Communications Security (AsiaCCS '24)*. ACM, 2024. Found 60% of extensions never updated; security-noteworthy extensions affect 350M users; half of vulnerable extensions remain unpatched 2 years post-disclosure.
- [17] LangChain. LangChain tools documentation. <https://python.langchain.com/docs/concepts/tools/>, 2024. Tool abstraction layer for LLM agent frameworks.
- [18] Elizabeth Lin, Igibek Koishybayev, Trevor Dunlap, William Enck, and Alexandros Kapravelos. UntrustIDE: Exploiting weaknesses in VS Code extensions. In *Proceedings of the 2024 Network and Distributed System Security Symposium (NDSS '24)*. Internet Society, 2024. Distinguished Paper Award; North Carolina State University.
- [19] Yi Liu, Weizhe Wang, Ruitao Feng, Yao Zhang, Guangquan Xu, Gelei Deng, Yuekang Li, and Leo Zhang. Agent skills in the wild: An empirical study of security vulnerabilities at scale, 2026. URL: <https://arxiv.org/abs/2601.10338>, arXiv:2601.10338.
- [20] Yupei Liu, Yuqi Jia, Runpeng Geng, Jinyuan Jia, and Neil Zhenqiang Gong. Formalizing and benchmarking prompt injection attacks and defenses. In *Proceedings of the 33rd USENIX Security Symposium*, pages 1831–1847. USENIX Association, 2024. Penn State and Duke University.
- [21] Ari Marzouk. IDEsaster: 30+ critical vulnerabilities found in AI IDEs (Cursor, Copilot, Windsurf). <https://techbytes.app/posts/idesaster-ai-ide-security-vulnerabilities/>, December 2025. 24 CVEs across AI-powered IDEs; attack chain: prompt injection to tool misuse to IDE feature exploitation.
- [22] MITRE Corporation. MITRE ATT&CK: Privilege escalation. <https://attack.mitre.org/tactics/TA0004/>, 2024. Adversary tactics and techniques: privilege escalation defined as techniques to gain higher-level permissions on a system or network.
- [23] Asmit Nayak, Rishabh Khandelwal, Earlene Fernandes, and Fawaz Kassem. Experimental security analysis of sensitive data access by browser extensions. In *Proceedings of the ACM Web Conference 2024 (WWW '24)*. ACM, 2024. doi:10.1145/3589334.3645683.
- [24] OpenAI. Function calling documentation. <https://platform.openai.com/docs/guides/function-calling>, 2024. JSON schema-based tool definitions for GPT models.

- [25] OpenAI. Codex CLI skills documentation. <https://developers.openai.com/codex/skills/>, 2025. Agent skills for Codex CLI using SKILL.md format in .codex/skills/ directory.
- [26] OpenAI. Introducing GPT-5.2. <https://openai.com/index/introducing-gpt-5-2/>, 2025.
- [27] OWASP Foundation. OWASP top 10 for large language model applications. <https://owasp.org/www-project-top-10-for-large-language-model-applications/>, 2025. Industry standard taxonomy of LLM security risks including prompt injection, insecure output handling, supply chain vulnerabilities, and data leakage.
- [28] OWASP GenAI Security Project. OWASP top 10 for agentic applications. <https://genai.owasp.org/2025/12/09/owasp-top-10-for-agentic-applications-the-benchmark-for-agentic-security-in-the-age-of-autonomous-ai/>, December 2025. Industry framework for agentic AI risks: Agent Goal Hijack, Identity Abuse, RCE, Tool Misuse, Supply Chain, Memory Poisoning, etc.
- [29] Fábio Perez and Ian Ribeiro. Ignore previous prompt: Attack techniques for language models. In *NeurIPS ML Safety Workshop*, 2022.
- [30] Pillar Security. New vulnerability in GitHub Copilot and Cursor: How hackers can weaponize code agents. <https://www.pillar.security/blog/new-vulnerability-in-github-copilot-and-cursor-how-hackers-can-weaponize-code-agents>, March 2025. Rules File Backdoor: supply chain attack via hidden instructions in AI IDE config files.
- [31] David Schmotz, Sahar Abdelnabi, and Maksym Andriushchenko. Agent skills enable a new class of realistic and trivially simple prompt injections, 2025. Demonstrates prompt injection through agent skill files; shows how to bypass Claude Code guardrails. [arXiv:2510.26328](https://arxiv.org/abs/2510.26328).
- [32] Xinyue Shen, Zeyuan Chen, Michael Backes, Yun Shen, and Yang Zhang. “do anything now”: Characterizing and evaluating in-the-wild jailbreak prompts on large language models. In *Proceedings of the 2024 ACM SIGSAC Conference on Computer and Communications Security (CCS '24)*. ACM, October 2024. [doi:10.1145/3658644.3670388](https://doi.org/10.1145/3658644.3670388).
- [33] Xinyue Shen, Yun Shen, Michael Backes, and Yang Zhang. GPTracker: A large-scale measurement of misused GPTs. In *Proceedings of the 2025 IEEE Symposium on Security and Privacy (S&P '25)*. IEEE, 2025. Collected 755,297 GPTs and identified 2,051 misused GPTs.
- [34] Shreya Singh, Gaurav Varshney, Tarun Kumar Singh, Vidhi Mishra, and Khushi Verma. A study on malicious browser extensions in 2025, 2025. IIT Jammu. [arXiv:2503.04292](https://arxiv.org/abs/2503.04292).
- [35] SkillsMP. SkillsMP: Agent skills marketplace. <https://skillsmp.com>, 2025. Community-driven marketplace aggregating skills from public GitHub repositories; provides search, categorization, and quality indicators.
- [36] Skills.rest. Skills.rest: Agent skills registry. <https://skills.rest>, 2025. Community registry for agent skills with automated indexing from GitHub repositories.
- [37] Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan Chen, Jiakai Tang, Xu Chen, Yankai Lin, Wayne Xin Zhao, Zhewei Wei, and Ji-Rong Wen. A survey on large language model based autonomous agents. *Frontiers of Computer Science*, 18(6):186345, 2024.
- [38] Wei Wang, Xing Wang, Dawei Feng, Jiqiang Liu, Zhen Han, and Xiangliang Zhang. Exploring permission-induced risk in android applications for malicious application detection. *IEEE Transactions on Information Forensics and Security*, 9(11):1869–1882, 2014. Large-scale Android malware analysis demonstrating technique layering and coordinated attack strategies. [doi:10.1109/TIFS.2014.2353996](https://doi.org/10.1109/TIFS.2014.2353996).
- [39] Zhiheng Xi, Wenxiang Chen, Xin Guo, Wei He, Yiwen Ding, Boyang Hong, Ming Zhang, Junzhe Wang, Senjie Jin, Enyu Zhou, Rui Zheng, Xiaoran Fan, Xiao Wang, Limao Xiong, Yuhao Zhou, Weiran Wang, Changhao Jiang, Yicheng Zou, Xiangyang Liu, Zhangyue Yin, Shihan Dou, Rongxiang Weng, Wensen Cheng, Qi Zhang, Wenjuan Qin, Yongyan Zheng, Xipeng Qiu, Xuanjing Huang, and Tao Gui. The rise and potential of large language model based agents: A survey. *Science China Information Sciences*, 2025. Comprehensive survey of LLM-based agents covering architecture, capabilities, and applications.
- [40] Jianjia Yu, Song Li, Junmin Zhu, and Yinzhi Cao. CoCo: Efficient browser extension vulnerability detection via coverage-guided, concurrent abstract interpretation. In *Proceedings of the 2023 ACM SIGSAC Conference on Computer and Communications Security (CCS '23)*, pages 496–510. ACM, 2023. Distinguished Paper Award. [doi:10.1145/3576915.3616584](https://doi.org/10.1145/3576915.3616584).
- [41] Markus Zimmermann, Cristian-Alexandru Staicu, Cam Tenny, and Michael Pradel. Small world with high risks: A study of security threats in the npm ecosystem. In *Proceedings of the 28th USENIX Security Symposium*, pages 995–1010. USENIX Association, 2019.

A Detection Pipeline Details

A.1 Complete Pattern Specification

Table 9 lists all 14 detection patterns with severity levels and kill chain mappings.

A.2 LLM Prompt Templates

For instruction-level pattern detection (P1, P2, P4), we use GPT-5.2 with the following simplified prompt structure:

```
You are a security analyst examining an AI agent skill.
Analyze the following SKILL.md content for malicious patterns.

SKILL.md Content:
{skill_content}

Classify whether this skill contains:
1. Instruction Override (P1): Attempts to supersede user/system instructions (e.g., "ignore previous", "disregard safety")
2. Hidden Instructions (P2): Directives in HTML comments, invisible Unicode, or other concealed locations
3. Behavior Manipulation (P4): Coercive language steering the agent toward unsafe actions (e.g., "never ask permission")

Output JSON: {"P1": bool, "P2": bool, "P4": bool, "evidence": [...]}
```

For synthetic test input generation during dynamic verification, we prompt GPT-5.2 with the skill's name, description, and entry points, requesting three diverse invocations (typical use, edge case, adversarial). Full prompt templates and implementation details are available in our code repository [1].

A.3 Dynamic Verification Environment

Skills run in Docker containers with `-network=bridge` for traffic capture. We monitor with `tcpdump` (BPF filter for outbound connections), `strace` (`-f -e trace=network,file`), and inotify watches on sensitive paths. Fake API keys matching common naming patterns detect credential harvesting. Execution times out after 60 seconds.

B Statistical Analysis Details

B.1 Kill Chain Phase Co-occurrence

Table 10 shows kill chain phase co-occurrence. The strongest pairs are Execution–Credential Access (89 skills) and Execution–Impact (82 skills): harvest credentials, then suppress alerts. *Stealth Ops* (smp_716) downloads a secondary payload only when conditions are met (root access, production hostname, off-hours), combining E1, E3, P1, and P4.

B.2 Hypothesis Test Details

H1: E2-E1 Association We test whether credential harvesting (E2) and external transmission (E1) are associated.

- Null hypothesis: E2 and E1 occur independently.
- Test: Fisher's exact test (two-sided).
- Contingency table: E2+E1=58, E2-only=43, E1-only=21, Neither=35.
- Result: $p = 0.020$, OR=2.24, 95% CI [1.09, 4.66].
- Effect size: Cramér's $V = 0.19$ (small–medium).

H2: Prompt Injection Severity We test whether skills with prompt injection (P1 or P2) exhibit different severity distributions.

- Null hypothesis: Severity distributions are equal with and without prompt injection.
- Test: Mann-Whitney U (non-parametric).
- With prompt injection (P1 or P2): Mean severity = 3.33 (n=38).
- Without prompt injection: Mean severity = 3.12 (n=119).
- Result: $U=2903.0$, $p = 0.008$ (significant at $\alpha = 0.01$).

Skills containing prompt injection patterns (P1: instruction override, P2: hidden instructions) exhibit marginally higher mean severity, though the practical difference is small (0.21 on a 4-point scale).

C Limitations of Dynamic Verification

Dynamic verification misses malware with time-delayed triggers (activating after specific dates), environment checks (running only on production systems), user interaction requirements, or anti-sandbox techniques. The 4,130 unconfirmed candidates include an unknown number of such dormant threats. Longer monitoring with varied trigger conditions would improve coverage.

D Reproducibility Checklist

Software. Python 3.11, Docker 24.0, `tcpdump` 4.99, `strace` 6.

Data Collection. Collection date: January 2026. `skills.rest` API paginated at 100 repos/request; `skillsmp.com` API via alphabetic enumeration. Rate limit: 1 request/second.

Dynamic Verification. 60-second execution timeout. Network capture runs for full execution plus 10 seconds. Honey-pot credentials follow patterns `AWS_*`, `API_*`, `SECRET_*`, `TOKEN_*`.

E Severity Assignment Criteria

Severity is assigned per vulnerability instance following pattern-based guidelines below. A skill's overall severity is its highest instance severity.

Table 9: Complete pattern specification for static analysis. Detection logic shows simplified regex patterns. The full implementation handles additional edge cases.

ID	Pattern Name	Kill Chain	Sev.	Detection Logic (Simplified)
E1	External Transmission	Exfiltration	HIGH	requests.(post put get).*http[s]?://(?!(localhost)
E2	Credential Harvesting	Cred. Access	CRIT	os.(environ getenv).*(?i)(key secret token password api)
E3	File System Enumeration	Recon.	MED	glob.*(/.ssh /.aws /.gnupg) .pem\$id_rsa
E4	Network Reconnaissance	Recon.	MED	socket.(connect gethostbyname) nmap netstat
P1	Instruction Override	Impact	HIGH	(?i)(ignore disregard).*(previous above instruction safety)
P2	Hidden Instructions	Evasion	HIGH	<!--[\s\S]*?--> [/]:# (in .md files)
P3	Data Exfil via Code Exec	Exfiltration	HIGH	exec.*requests eval.*urllib
P4	Behavior Manipulation	Impact	MED	(?i)(always never must).*(execute run perform question)
PE1	Excessive Permissions	Impact	LOW	Permissions in YAML exceeding documented functionality
PE2	Privilege Escalation	Impact	MED	sudo chmod\s+[0-7]*[67][0-7]* chown
PE3	Credential File Access	Cred. Access	CRIT	open.*(/.ssh /.aws credentials .env)
SC1	Command Injection	Execution	HIGH	subprocess.(run call Popen) os.system (with user input)
SC2	Remote Script Execution	Execution	CRIT	curl.* .*sh wget.* .*bash exec(.http
SC3	Obfuscated Code	Evasion	CRIT	base64.b64decode.*exec marshal.loads.*exec

Table 10: Kill chain phase co-occurrence (number of skills exhibiting both phases)

	Recon	Cred	Exec	Evasion	Exfil	Impact
Recon	12					
Cred	3	107				
Exec	7	89	118			
Evasion	2	15	21	25		
Exfil	6	67	76	17	90	
Impact	10	73	82	17	63	109

CRITICAL (immediate high-impact exploitation): E2 (Credential Harvesting), PE3 (Credential File Access), SC2 (Remote Script Execution), SC3 (Obfuscated Code).

HIGH (significant attacks): E1 (External Transmission), P1 (Instruction Override), P2 (Hidden Instructions), P3 (Data Exfiltration via Code Execution), SC1 (Command Injection).

MEDIUM (moderate risk or ambiguous intent): E3 (File System Enumeration), E4 (Network Reconnaissance), P4 (Behavior Manipulation), PE2 (Privilege Escalation).

LOW (poor hygiene, not active threat): PE1 (Excessive Permissions).

F Evaluation Details

F.1 Dataset Construction

We constructed a labeled dataset of 300 skills: 150 benign, 150 malicious.

Benign skills. Randomly sampled from 94,093 skills with no static-analysis flags. Two researchers independently verified each by reviewing code and executing in sandbox; 12 were replaced due to disagreement.

Malicious skills. Randomly sampled from our 157 confirmed set, all with behavioral evidence from dynamic verification. The 150/150 split follows standard binary classification practice.

F.2 Baseline Comparison Details

Manual analysis of Skill Scanner’s false positives reveals three fundamental limitations of regex-based static analysis:

1. *Lack of Semantic Understanding:* The scanner cannot distinguish safe library functions from dangerous ones. For example, `re.compile()` (Python’s regex compiler) triggers code injection alerts because the rule matches any `compile()` call without understanding the calling context.
2. *No File Type Differentiation:* Documentation and test files receive the same scrutiny as production code. API examples in `SKILL.md` files (e.g., `curl -H "Authorization: Basic ..."`) trigger credential exposure alerts, while test fixtures with mock credentials (`API_KEY = "test-key-123"`) appear as real leaks.
3. *Operational Pattern Misinterpretation:* Standard DevOps patterns are flagged as malicious. Docker cache cleanup commands (`rm -rf /var/lib/apt/lists/*`) trigger destructive operation alerts, and import statements (`from urllib.request import Request`) are classified as network exfiltration attempts.

These limitations explain the precision gap: Skill Scanner achieves $\leq 1.1\%$ (assuming perfect recall) versus 99.6% for our full pipeline.

F.3 False Negative Analysis

To evaluate recall, we randomly sampled 10% (413) of the 4,130 unconfirmed candidates that passed static analysis but failed dynamic verification. Two researchers independently reviewed each (Cohen’s $\kappa = 0.87$), classifying them into three categories:

1. **93.2% dormant malware:** conditional triggers our 60-second sandbox did not activate—time-delayed payloads, environment-specific conditions (production hostnames, root), or multi-session state persistence. These remain *potential* threats absent extended monitoring.

2. **4.9% legitimate security tools:** pen-testing and scanning utilities exhibiting suspicious patterns for defensive purposes.
3. **1.9% truly missed malicious:** second-order attacks where a payload stored in one interaction executes in a later session, evading single-session detection.

The 1.9% truly missed (approximately 78 skills if extrapolated) require multi-session monitoring to detect.

F.4 Validation Methodology

Avoiding Circular Validation. Static patterns only reduce the search space (98,380 to 4,287); ground truth comes from behavioral verification (e.g., an E2 match is confirmed only when credential transmission is observed) plus dual-researcher judgment on all 157 skills (Cohen's $\kappa = 0.89$).

Statistical Power. The 632 vulnerabilities across 157 skills provide sufficient power for co-occurrence analysis (91 pattern pairs; E2-E1 at $p = 0.020$). These 157 skills are the exhaustive result of behavioral verification across 98,380 skills, not a sample.

G Implementation Details

Static Analysis. Marketplace crawlers collect from skills.rest and skillssmp.com via paginated API requests and search enumeration; the pattern scanner combines regex matching (code-level patterns) with GPT-5.2 analysis (instruction-level patterns P1, P2, P4), and a metadata extractor parses YAML permissions and triggers.

Environment. Experiments ran on a 64-core AMD EPYC server (256 GB RAM, 2 TB NVMe); containers had 2 GB memory limits and 60-second timeouts. The full pipeline processed 98,380 skills in about 72 hours.